



JOINT INSTITUTE
交大密西根学院

VE489

COMPUTER NETWORK

REPORT OF PROJECT

2020 SUMMER

Name:

JUNJIE ZHANG

Student Number:

517021911128

Instructor:

DR. XUDONG WANG

JULY 28, 2020

Contents

1	Part I	2
1.1	Link latency using ping	2
1.2	Path latency using ping	2
1.3	Link bandwidth using iperf	3
1.4	Path throughput using iperf	4
1.5	Multiplexing	5
2	Part II	7
2.1	Implementation	7
2.2	Results	8
2.2.1	Transmits file with size of 100B	8
2.2.2	Transmits file with size of 100KB	9
2.2.3	Transmits file with size of 10MB	9
3	Part III	12
3.1	Implement a simple SR	12
3.1.1	Implementation	12
3.1.2	Test and Results	13
3.2	Make sender more efficient – leverage duplicate ACK on sender side	18
3.2.1	Implementation	18
3.2.2	Tests and Results	18
3.3	Make sender more efficient – send NACK on receiver side	18
3.3.1	Implementation	18
3.3.2	Tests and Results	18
3.4	Throughput, delay and window size	19
4	Appendix	19
4.1	Part II	19
4.1.1	makeFile	19
4.1.2	ftrans.c	19
4.2	Part III	20
4.2.1	makeFile	20
4.2.2	sr.c in 3.1	20
4.2.3	sr.c in 3.2	28
4.2.4	sr.c in 3.3	37

1 Part I

1.1 Link latency using ping

```
mininet> h1 ping -c 10 h2
PING 10.0.0.2 (10.0.0.2) 56(84) bytes of data.
64 bytes from 10.0.0.2: icmp_seq=1 ttl=64 time=44.2 ms
64 bytes from 10.0.0.2: icmp_seq=2 ttl=64 time=43.8 ms
64 bytes from 10.0.0.2: icmp_seq=3 ttl=64 time=41.7 ms
64 bytes from 10.0.0.2: icmp_seq=4 ttl=64 time=41.3 ms
64 bytes from 10.0.0.2: icmp_seq=5 ttl=64 time=41.4 ms
64 bytes from 10.0.0.2: icmp_seq=6 ttl=64 time=40.0 ms
64 bytes from 10.0.0.2: icmp_seq=7 ttl=64 time=40.0 ms
64 bytes from 10.0.0.2: icmp_seq=8 ttl=64 time=40.6 ms
64 bytes from 10.0.0.2: icmp_seq=9 ttl=64 time=40.0 ms
64 bytes from 10.0.0.2: icmp_seq=10 ttl=64 time=40.0 ms

--- 10.0.0.2 ping statistics ---
10 packets transmitted, 10 received, 0% packet loss, time 9017ms
rtt min/avg/max/mdev = 40.042/41.349/44.251/1.506 ms
```

Figure 1: Ping result between $h1$ and $h2$

From Figure 1 we can see that the average RTT is $41.348ms$ between $h1$ and $h2$, which indicates $L1$ link is used, and RTT is close to 2 times link latency.

```
mininet> h3 ping -c 10 h5
PING 10.0.0.5 (10.0.0.5) 56(84) bytes of data.
64 bytes from 10.0.0.5: icmp_seq=1 ttl=64 time=23.4 ms
64 bytes from 10.0.0.5: icmp_seq=2 ttl=64 time=22.7 ms
64 bytes from 10.0.0.5: icmp_seq=3 ttl=64 time=21.5 ms
64 bytes from 10.0.0.5: icmp_seq=4 ttl=64 time=20.8 ms
64 bytes from 10.0.0.5: icmp_seq=5 ttl=64 time=20.6 ms
64 bytes from 10.0.0.5: icmp_seq=6 ttl=64 time=21.2 ms
64 bytes from 10.0.0.5: icmp_seq=7 ttl=64 time=20.7 ms
64 bytes from 10.0.0.5: icmp_seq=8 ttl=64 time=20.6 ms
64 bytes from 10.0.0.5: icmp_seq=9 ttl=64 time=21.5 ms
64 bytes from 10.0.0.5: icmp_seq=10 ttl=64 time=20.5 ms

--- 10.0.0.5 ping statistics ---
10 packets transmitted, 10 received, 0% packet loss, time 9012ms
rtt min/avg/max/mdev = 20.571/21.392/23.434/0.940 ms
```

Figure 2: Ping result between $h3$ and $h5$

From Figure 2 we can see that the average RTT is $21.392ms$ between $h3$ and $h5$, which indicates $L1$ link is used, and RTT is close to 2 times link latency.

1.2 Path latency using ping

For $h1$ and $h5$, the path should be $L1 \leftrightarrow L2 \leftrightarrow L4$, then the theoretical RTT is:

$$RTT_{(h1,h5)the} = 2 \times (40 + 20 + 10) = 140ms$$

From Figure 3, we can see that the average RTT is $142.530ms$, which is very close to the theoretical one.

```

mininet> h1 ping -c 10 h5
PING 10.0.0.5 (10.0.0.5) 56(84) bytes of data.
64 bytes from 10.0.0.5: icmp_seq=1 ttl=64 time=147 ms
64 bytes from 10.0.0.5: icmp_seq=2 ttl=64 time=146 ms
64 bytes from 10.0.0.5: icmp_seq=3 ttl=64 time=143 ms
64 bytes from 10.0.0.5: icmp_seq=4 ttl=64 time=141 ms
64 bytes from 10.0.0.5: icmp_seq=5 ttl=64 time=141 ms
64 bytes from 10.0.0.5: icmp_seq=6 ttl=64 time=140 ms
64 bytes from 10.0.0.5: icmp_seq=7 ttl=64 time=140 ms
64 bytes from 10.0.0.5: icmp_seq=8 ttl=64 time=141 ms
64 bytes from 10.0.0.5: icmp_seq=9 ttl=64 time=140 ms
64 bytes from 10.0.0.5: icmp_seq=10 ttl=64 time=141 ms

--- 10.0.0.5 ping statistics ---
10 packets transmitted, 10 received, 0% packet loss, time 9013ms
rtt min/avg/max/mdev = 140.100/142.530/147.774/2.576 ms

```

Figure 3: Ping result between $h1$ and $h5$

For $h3$ and $h4$, the path should be $L2 \leftrightarrow L3$, then the theoretical RTT is:

$$RTT_{(h3,h4)the} = 2 \times (40 + 30) = 140ms$$

From Figure 4, we can see that the average RTT is $142.700ms$, which is very close to the theoretical one.

```

mininet> h3 ping -c 10 h4
PING 10.0.0.4 (10.0.0.4) 56(84) bytes of data.
64 bytes from 10.0.0.4: icmp_seq=1 ttl=64 time=147 ms
64 bytes from 10.0.0.4: icmp_seq=2 ttl=64 time=146 ms
64 bytes from 10.0.0.4: icmp_seq=3 ttl=64 time=142 ms
64 bytes from 10.0.0.4: icmp_seq=4 ttl=64 time=141 ms
64 bytes from 10.0.0.4: icmp_seq=5 ttl=64 time=141 ms
64 bytes from 10.0.0.4: icmp_seq=6 ttl=64 time=142 ms
64 bytes from 10.0.0.4: icmp_seq=7 ttl=64 time=140 ms
64 bytes from 10.0.0.4: icmp_seq=8 ttl=64 time=140 ms
64 bytes from 10.0.0.4: icmp_seq=9 ttl=64 time=141 ms
64 bytes from 10.0.0.4: icmp_seq=10 ttl=64 time=142 ms

--- 10.0.0.4 ping statistics ---
10 packets transmitted, 10 received, 0% packet loss, time 9016ms
rtt min/avg/max/mdev = 140.626/142.700/147.103/2.185 ms

```

Figure 4: Ping result between $h3$ and $h4$

1.3 Link bandwidth using iperf

For $h1$ and $h2$, let $h1$ be the server and let $h2$ be the client. Then the bandwidth of $h1$ can be seen in Figure 5, which is $46.8Mbits/s$. And the bandwidth of $h2$ can be seen in Figure 6, which is $56.3Mbits/s$. The bandwidth of $L1$ is $50Mbits/s$, which is close to above two. And from Figure 5 and Figure 6 we can see that the sizes of data transferred and received are the same, which is $35.6MBytes$.

For $h3$ and $h5$, let $h3$ be the server and let $h5$ be the client. Then the bandwidth of $h3$ can be seen in Figure 7, which is $9.55Mbits/s$. And the bandwidth of $h2$ can be seen in Figure 8, which is $11.5Mbits/s$. The bandwidth of $L1$ is $10Mbits/s$, which is close to above two. And from Figure 7 and

```

root@ubuntu:~/Project/part1# iperf -s -p 8080
-----
Server listening on TCP port 8080
TCP window size: 85.3 KByte (default)
-----
[ 28] local 10.0.0.1 port 8080 connected with 10.0.0.2 port 33219
[ ID] Interval      Transfer    Bandwidth
[ 28] 0.0- 6.4 sec   35.6 MBytes 46.8 Mbits/sec

```

Figure 5: Server on $h1$

```

root@ubuntu:~/Project/part1# iperf -c 10.0.0.1 -p 8080 -t 5
-----
Client connecting to 10.0.0.1, TCP port 8080
TCP window size: 85.3 KByte (default)
-----
[ 27] local 10.0.0.2 port 33219 connected with 10.0.0.1 port 8080
[ ID] Interval      Transfer    Bandwidth
[ 27] 0.0- 5.3 sec   35.6 MBytes 56.3 Mbits/sec

```

Figure 6: Client on $h2$

Figure 8 we can see that the sizes of data transferred and received are the same, which is $7.00MBytes$.

```

root@ubuntu:~/Project/part1# iperf -s -p 8080
-----
Server listening on TCP port 8080
TCP window size: 85.3 KByte (default)
-----
[ 28] local 10.0.0.3 port 8080 connected with 10.0.0.5 port 37502
[ ID] Interval      Transfer    Bandwidth
[ 28] 0.0- 6.1 sec    7.00 MBytes 9.55 Mbits/sec

```

Figure 7: Server on $h3$

```

root@ubuntu:~/Project/part1# iperf -c 10.0.0.3 -p 8080 -t 5
-----
Client connecting to 10.0.0.3, TCP port 8080
TCP window size: 85.3 KByte (default)
-----
[ 27] local 10.0.0.5 port 37502 connected with 10.0.0.3 port 8080
[ ID] Interval      Transfer    Bandwidth
[ 27] 0.0- 5.1 sec    7.00 MBytes 11.5 Mbits/sec

```

Figure 8: Client on $h5$

1.4 Path throughput using iperf

For $h1$ and $h5$, let $h1$ be the server and let $h5$ be the client. Then the bandwidth of $h1$ can be seen in Figure 9, which is $8.84Mbits/s$. And the bandwidth of $h5$ can be seen in Figure 10, which is $11.1Mbits/s$. The bottleneck link is $L4$, whose bandwidth is $10Mbits/s$, which is close to above two. And from Figure 9 and Figure 10 we can see that the sizes of data transferred and received are the same, which is $7.25MBytes$.

For $h3$ and $h4$, let $h3$ be the server and let $h4$ be the client. Then the bandwidth of $h3$ can be seen in Figure 11, which is $17.3Mbits/s$. And the bandwidth of $h5$ can be seen in Figure 12, which is

```

root@ubuntu:~/Project/part1# iperf -s -p 8080
-----
Server listening on TCP port 8080
TCP window size: 85.3 KByte (default)
-----
[ 28] local 10.0.0.1 port 8080 connected with 10.0.0.5 port 45934
[ ID] Interval      Transfer    Bandwidth
[ 28] 0.0- 6.9 sec  7.25 MBytes  8.84 Mbits/sec

```

Figure 9: Server on *h1*

```

root@ubuntu:~/Project/part1# iperf -c 10.0.0.1 -p 8080 -t 5
-----
Client connecting to 10.0.0.1, TCP port 8080
TCP window size: 85.3 KByte (default)
-----
[ 27] local 10.0.0.5 port 45934 connected with 10.0.0.1 port 8080
[ ID] Interval      Transfer    Bandwidth
[ 27] 0.0- 5.5 sec  7.25 MBytes  11.1 Mbits/sec

```

Figure 10: Client on *h5*

19.1Mbits/s. The bottleneck link is *L3*, whose bandwidth is 20Mbits/s, which is close to above two. And from Figure 11 and Figure 12 we can see that the sizes of data transferred and received are the same, which is 11.5MBytes.

```

root@ubuntu:~/Project/part1# iperf -s -p 8080
-----
Server listening on TCP port 8080
TCP window size: 85.3 KByte (default)
-----
[ 28] local 10.0.0.3 port 8080 connected with 10.0.0.4 port 44441
[ ID] Interval      Transfer    Bandwidth
[ 28] 0.0- 5.6 sec  11.5 MBytes  17.3 Mbits/sec

```

Figure 11: Server on *h3*

```

root@ubuntu:~/Project/part1# iperf -c 10.0.0.3 -p 8080 -t 5
-----
Client connecting to 10.0.0.3, TCP port 8080
TCP window size: 85.3 KByte (default)
-----
[ 27] local 10.0.0.4 port 44441 connected with 10.0.0.3 port 8080
[ ID] Interval      Transfer    Bandwidth
[ 27] 0.0- 5.0 sec  11.5 MBytes  19.1 Mbits/sec

```

Figure 12: Client on *h4*

1.5 Multiplexing

For the latency. From Figure 13, we can see that the average *RTT* between *h1* and *h5* is 141.702ms, which is close to the theoretical one. And from Figure 14, we can see that the average *RTT* between *h3* and 45 is 142.995ms, which is also close to the theoretical one. Compared with question 2, the result stay unchanged. And link *L2* is shared, since the rate of sending packet by ping is much less than the bandwidth of *L2*, we can say that the link is shared fairly and the result won't change.

```

root@ubuntu:~/Project/part1# ping -c 20 10.0.0.4
PING 10.0.0.4 (10.0.0.4) 56(84) bytes of data.
64 bytes from 10.0.0.4: icmp_seq=1 ttl=64 time=145 ms
64 bytes from 10.0.0.4: icmp_seq=2 ttl=64 time=141 ms
64 bytes from 10.0.0.4: icmp_seq=3 ttl=64 time=141 ms
64 bytes from 10.0.0.4: icmp_seq=4 ttl=64 time=140 ms
64 bytes from 10.0.0.4: icmp_seq=5 ttl=64 time=141 ms
64 bytes from 10.0.0.4: icmp_seq=6 ttl=64 time=141 ms
64 bytes from 10.0.0.4: icmp_seq=7 ttl=64 time=140 ms
64 bytes from 10.0.0.4: icmp_seq=8 ttl=64 time=141 ms
64 bytes from 10.0.0.4: icmp_seq=9 ttl=64 time=141 ms
64 bytes from 10.0.0.4: icmp_seq=10 ttl=64 time=142 ms
64 bytes from 10.0.0.4: icmp_seq=11 ttl=64 time=142 ms
64 bytes from 10.0.0.4: icmp_seq=12 ttl=64 time=141 ms
64 bytes from 10.0.0.4: icmp_seq=13 ttl=64 time=140 ms
64 bytes from 10.0.0.4: icmp_seq=14 ttl=64 time=141 ms
64 bytes from 10.0.0.4: icmp_seq=15 ttl=64 time=141 ms
64 bytes from 10.0.0.4: icmp_seq=16 ttl=64 time=141 ms
64 bytes from 10.0.0.4: icmp_seq=17 ttl=64 time=141 ms
64 bytes from 10.0.0.4: icmp_seq=18 ttl=64 time=141 ms
64 bytes from 10.0.0.4: icmp_seq=19 ttl=64 time=141 ms
64 bytes from 10.0.0.4: icmp_seq=20 ttl=64 time=142 ms

--- 10.0.0.4 ping statistics ---
20 packets transmitted, 20 received, 0% packet loss, time 19025ms
rtt min/avg/max/mdev = 140.600/141.702/145.988/1.116 ms

```

Figure 13: Ping result between $h1$ and $h5$

```

root@ubuntu:~/Project/part1# ping -c 20 10.0.0.5
PING 10.0.0.5 (10.0.0.5) 56(84) bytes of data.
64 bytes from 10.0.0.5: icmp_seq=1 ttl=64 time=145 ms
64 bytes from 10.0.0.5: icmp_seq=2 ttl=64 time=142 ms
64 bytes from 10.0.0.5: icmp_seq=3 ttl=64 time=143 ms
64 bytes from 10.0.0.5: icmp_seq=4 ttl=64 time=143 ms
64 bytes from 10.0.0.5: icmp_seq=5 ttl=64 time=142 ms
64 bytes from 10.0.0.5: icmp_seq=6 ttl=64 time=143 ms
64 bytes from 10.0.0.5: icmp_seq=7 ttl=64 time=142 ms
64 bytes from 10.0.0.5: icmp_seq=8 ttl=64 time=142 ms
64 bytes from 10.0.0.5: icmp_seq=9 ttl=64 time=142 ms
64 bytes from 10.0.0.5: icmp_seq=10 ttl=64 time=141 ms
64 bytes from 10.0.0.5: icmp_seq=11 ttl=64 time=143 ms
64 bytes from 10.0.0.5: icmp_seq=12 ttl=64 time=142 ms
64 bytes from 10.0.0.5: icmp_seq=13 ttl=64 time=141 ms
64 bytes from 10.0.0.5: icmp_seq=14 ttl=64 time=144 ms
64 bytes from 10.0.0.5: icmp_seq=15 ttl=64 time=144 ms
64 bytes from 10.0.0.5: icmp_seq=16 ttl=64 time=142 ms
64 bytes from 10.0.0.5: icmp_seq=17 ttl=64 time=142 ms
64 bytes from 10.0.0.5: icmp_seq=18 ttl=64 time=143 ms
64 bytes from 10.0.0.5: icmp_seq=19 ttl=64 time=142 ms
64 bytes from 10.0.0.5: icmp_seq=20 ttl=64 time=144 ms

--- 10.0.0.5 ping statistics ---
20 packets transmitted, 20 received, 0% packet loss, time 19009ms
rtt min/avg/max/mdev = 141.826/142.995/145.139/1.056 ms

```

Figure 14: Ping result between $h1$ and $h5$

For the bandwidth. Let $h1$ and $h3$ be servers. Let $h3$ and $h4$ be clients. From Figure 15, we can see that the bandwidth for $h1$ is 2.33Mbits/s and the bandwidth for $h5$ is 2.33Mbits/s . Both of them are not close to the bandwidth of the bottleneck link, $L4$, which is 10Mbits/s . From Figure 16, we can see that the bandwidth for $h1$ is 17.6Mbits/s and the bandwidth for $h5$ is 22.8Mbits/s . Both of them are relatively close to the bandwidth of the bottleneck link, $L2$, which is 20Mbits/s . This indicate that the result is not the same as the one of question 4. The shared link is $L2$ whose bandwidth is 20Mbits/s . Since the total bandwidth of two bottleneck link is 30Mbits , which is larger than the bandwidth of the shared link, it means the result will differ from the individual one. And clearly, $L2$ is not shared fairly.


```
"Node: h1"
root@ubuntu:~/Project/part1# iperf -s -p 8080
-----
Server listening on TCP port 8080
TCP window size: 85.3 KByte (default)
-----
[ 28] local 10.0.0.1 port 8080 connected with 10.0.0.5 port 45965
[ ID] Interval      Transfer    Bandwidth
[ 28] 0.0- 6.1 sec  1.62 MBytes  2.22 Mbits/sec

"Node: h5"
root@ubuntu:~/Project/part1# iperf -c 10.0.0.1 -p 8080 -t 5
-----
Client connecting to 10.0.0.1, TCP port 8080
TCP window size: 85.3 KByte (default)
-----
[ 27] local 10.0.0.5 port 45965 connected with 10.0.0.1 port 8080
[ ID] Interval      Transfer    Bandwidth
[ 27] 0.0- 5.8 sec  1.62 MBytes  2.33 Mbits/sec
```

Figure 15: Server on *h1* and Client on *h5*

```
"Node: h3"
root@ubuntu:~/Project/part1# iperf -s -p 8080
-----
Server listening on TCP port 8080
TCP window size: 85.3 KByte (default)
-----
[ 28] local 10.0.0.3 port 8080 connected with 10.0.0.4 port 44470
[ ID] Interval      Transfer    Bandwidth
[ 28] 0.0- 6.7 sec  14.0 MBytes  17.6 Mbits/sec

"Node: h4"
root@ubuntu:~/Project/part1# iperf -c 10.0.0.3 -p 8080 -t 5
-----
Client connecting to 10.0.0.3, TCP port 8080
TCP window size: 85.3 KByte (default)
-----
[ 27] local 10.0.0.4 port 44470 connected with 10.0.0.3 port 8080
[ ID] Interval      Transfer    Bandwidth
[ 27] 0.0- 5.2 sec  14.0 MBytes  22.8 Mbits/sec
```

Figure 16: Server on *h3* and Client on *h4*

2 Part II

2.1 Implementation

The whole process is as follows:

1. Client sends *filename* to server.
2. Server reads *filename* and open file.
3. Server sends *length* to client.
4. Client reads *length*.
5. Server sends *file_data* to client.
6. Client reads *file_data* (length bytes in total) and writes to *filename*.

No.	Time	Source	Destination	Protocol	Length	Info
1	0.000000	0.0.0.0	255.255.255.255	DHCP	342	DHCP Discover - Transaction ID 0x874d5c57
2	0.635346	0.0.0.0	255.255.255.255	DHCP	342	DHCP Discover - Transaction ID 0x2388b03f
3	1.405294	0.0.0.0	255.255.255.255	DHCP	342	DHCP Discover - Transaction ID 0xeb0da147
4	1.926963	0.0.0.0	255.255.255.255	DHCP	342	DHCP Discover - Transaction ID 0xf6ee642
5	4.711673	10.0.0.2	10.0.0.1	TCP	74	34433 → 1000 [SYN] Seq=0 Win=29200 Len=0 MSS=1460 SACK_PERM=1 TSval=8915888 TSecr=0 WS=512
6	4.755167	10.0.0.1	10.0.0.2	TCP	74	1000 → 34433 [SYN, ACK] Seq=0 Ack=1 Win=28960 Len=0 MSS=1460 SACK_PERM=1 TSval=8915894 TSecr=8915888 WS=512
7	4.755188	10.0.0.2	10.0.0.1	TCP	66	34433 → 1000 [ACK] Seq=1 Ack=1 Win=29696 Len=0 TSval=8915899 TSecr=8915894
8	4.756358	10.0.0.2	10.0.0.1	TCP	71	34433 → 1000 [PSH, ACK] Seq=1 Ack=1 Win=29696 Len=5 TSval=8915899 TSecr=8915894
9	4.799622	10.0.0.1	10.0.0.2	TCP	66	1000 → 34433 [ACK] Seq=1 Ack=6 Win=29184 Len=0 TSval=8915905 TSecr=8915899
10	4.799649	10.0.0.1	10.0.0.2	TCP	74	1000 → 34433 [PSH, ACK] Seq=1 Ack=6 Win=29184 Len=8 TSval=8915905 TSecr=8915899
11	4.799657	10.0.0.2	10.0.0.1	TCP	66	34433 → 1000 [ACK] Seq=6 Ack=9 Win=29696 Len=0 TSval=8915910 TSecr=8915905
12	5.059118	10.0.0.1	10.0.0.2	TCP	74	[TCP Spurious Retransmission] 1000 → 34433 [PSH, ACK] Seq=1 Ack=6 Win=29184 Len=8 TSval=8915970 TSecr=8915899
13	5.059145	10.0.0.2	10.0.0.1	TCP	78	[TCP Dup ACK 11#1] 34433 → 1000 [ACK] Seq=6 Ack=9 Win=29696 Len=0 TSval=8915975 TSecr=8915975
14	5.802151	10.0.0.1	10.0.0.2	TCP	166	1000 → 34433 [PSH, ACK] Seq=9 Ack=6 Win=29184 Len=100 TSval=8916156 TSecr=8915975
15	5.802282	10.0.0.2	10.0.0.1	TCP	66	34433 → 1000 [ACK] Seq=6 Ack=109 Win=29696 Len=0 TSval=8916161 TSecr=8916156
16	5.803779	10.0.0.2	10.0.0.1	TCP	66	34433 → 1000 [FIN, ACK] Seq=6 Ack=109 Win=29696 Len=0 TSval=8916161 TSecr=8916156
17	5.803305	10.0.0.1	10.0.0.2	TCP	66	1000 → 34433 [ACK] Seq=109 Ack=7 Win=29184 Len=0 TSval=8916176 TSecr=8916161
18	6.803549	10.0.0.1	10.0.0.2	TCP	66	1000 → 34433 [FIN, ACK] Seq=109 Ack=7 Win=29184 Len=0 TSval=8916406 TSecr=8916161
19	6.803585	10.0.0.2	10.0.0.1	TCP	66	34433 → 1000 [ACK] Seq=7 Ack=110 Win=29696 Len=0 TSval=8916411 TSecr=8916406
20	7.106497	0.0.0.0	255.255.255.255	DHCP	342	DHCP Discover - Transaction ID 0xf6b455c
21	7.893967	fe80::1c1a:bcff:fe57:dff02::fb	ff02::fb	MDNS	192	Standard query 0x0000 PTR _afpovertcp._tcp.local, "QM" question PTR _ipp._tcp.local, "QM" question PTR _ipps._tcp.local.

Figure 17: The packets captured when transmit file with size of 100B

Frame 8: 71 bytes on wire (568 bits), 71 bytes captured (568 bits)
Ethernet II, Src: 00:00:00_00:00:02 (00:00:00:00:00:02), Dst: 00:00:00_00:00:01 (00:00:00:00:00:01)
Internet Protocol Version 4, Src: 10.0.0.2, Dst: 10.0.0.1
Transmission Control Protocol, Src Port: 34433, Dst Port: 1000, Seq: 1, Ack: 1, Len: 5
Data (5 bytes)
Data: 7465737431
[Length: 5]
0010 00 39 20 53 40 00 40 06 06 6a 0a 00 00 02 0a 00 .9 S@.@. .j.....
0020 00 01 86 81 03 e8 70 a3 2b e2 25 73 d3 6e 80 18p. +.%S.n..
0030 00 3a 14 2e 00 00 01 01 08 0a 00 88 0b bb 00 88
0040 0b b6 74 65 73 74 31 ..test1

Figure 18: Detailed information about packet No.8

7. Client exits, server keeps running.

And I choose 1024 bytes as the buffer size, the code of `ftrans.cc` and `Makefile` can be seen in **Appendix**.

2.2 Results

The server is set on *h1* and the client is set on *h2*.

2.2.1 Transmits file with size of 100B

The file is generated by command `dd if=/dev/zero of=test1 bs=1 count=100`

We captured the packets sent and received on client with `tcpdump`, and the result is shown through `Wireshark`, which can be seen in Figure 17.

Then, we can see that packet No.8 contains the file name, which is from *h2*(10.0.0.2) to *h1*(10.0.0.1) and can be seen in Figure 19. The file name is *test1*, which is correct.

And packet No.10 contains the file length, which is from *h1* to *h2*. The file length is $0x64 = 100$ bytes, which is correct.

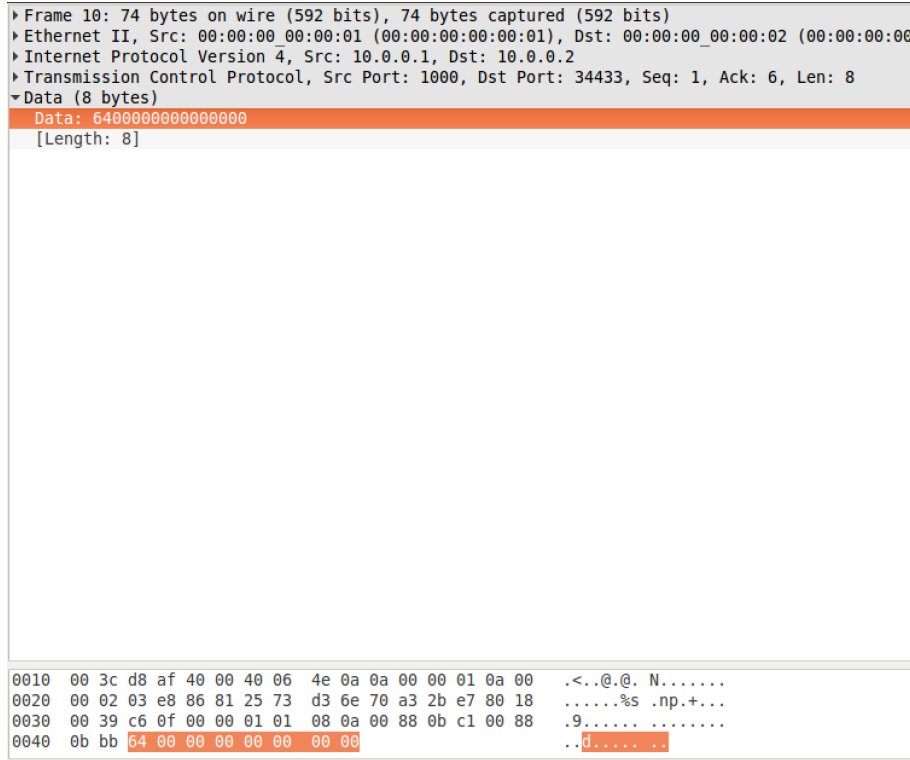


Figure 19: Detailed information about packet No.10

And packet No.14 contains the data, which is from $h1$ to $h2$. By using `diff`, we can say that the transmission is correct.

2.2.2 Transmits file with size of 100KB

The file is generated by command `dd if=/dev/zero of=test1 bs=1KB count=100`

We captured the packets sent and received on client with `tcpdump`, and the part of the result is shown through `Wireshark`, which can be seen in Figure 20.

Then, we can see that packet No.8 contains the file name, which is from $h2(10.0.0.2)$ to $h1(10.0.0.1)$ and can be seen in Figure 19. The file name is `test2`, which is correct.

And packet No.10 contains the file length, which is from $h1$ to $h2$. The file length is `0x0186a0 = 100000 bytes = 100KB`, which is correct.

And packets with size of more than 1000 bytes contain the data, which is from $h1$ to $h2$. By using `diff`, we can say that the transmission is correct.

2.2.3 Transmits file with size of 10MB

The file is generated by command `dd if=/dev/zero of=test1 bs=1MB count=10`

We captured the packets sent and received on client with `tcpdump`, and the part of the result is shown through `Wireshark`, which can be seen in Figure 23.

Then, we can see that packet No.22 contains the file name, which is from $h2(10.0.0.2)$ to $h1(10.0.0.1)$ and can be seen in Figure 24. The file name is `test3`, which is correct.

No.	Time	Source	Destination	Protocol	Length	Info
1	0.000000	0.0.0.0	255.255.255.255	DHCP	342	DHCP Discover - Transaction ID 0xeb0da147
2	0.618343	0.0.0.0	255.255.255.255	DHCP	342	DHCP Discover - Transaction ID 0x874d5c57
3	1.702542	0.0.0.0	255.255.255.255	DHCP	342	DHCP Discover - Transaction ID 0x2388b03f
4	4.629979	10.0.0.2	10.0.0.1	TCP	74	33228 → 1000 [SYN] Seq=0 Win=29200 Len=0 MSS=1460 SACK_PERM=1 TSval=8919049 TSecr=0 WS=512
5	5.630011	10.0.0.2	10.0.0.1	TCP	74	[TCP Retransmission] 33228 → 1000 [SYN] Seq=0 Win=29200 Len=0 MSS=1460 SACK_PERM=1 TSval=8919299 TSecr=0
6	5.673962	10.0.0.1	10.0.0.2	TCP	74	1000 → 33228 [SYN, ACK] Seq=0 Ack=1 Win=28960 Len=0 MSS=1460 SACK_PERM=1 TSval=8919304 TSecr=8919049
7	5.673983	10.0.0.2	10.0.0.1	TCP	66	33228 → 1000 [ACK] Seq=1 Ack=1 Win=29696 Len=0 TSval=8919310 TSecr=8919304
8	5.674276	10.0.0.2	10.0.0.1	TCP	71	33228 → 1000 [PSH, ACK] Seq=1 Ack=1 Win=29696 Len=5 TSval=8919310 TSecr=8919304
9	5.714709	10.0.0.1	10.0.0.2	TCP	66	1000 → 33228 [ACK] Seq=1 Ack=6 Win=29184 Len=0 TSval=8919315 TSecr=8919310
10	5.715228	10.0.0.1	10.0.0.2	TCP	74	1000 → 33228 [PSH, ACK] Seq=1 Ack=6 Win=29184 Len=8 TSval=8919315 TSecr=8919310
11	5.715240	10.0.0.2	10.0.0.1	TCP	66	33228 → 1000 [ACK] Seq=6 Ack=9 Win=29696 Len=0 TSval=8919320 TSecr=8919315
12	5.978710	10.0.0.1	10.0.0.2	TCP	74	[TCP Spurious Retransmission] 1000 → 33228 [PSH, ACK] Seq=1 Ack=6 Win=29184 Len=8 TSval=8919381 TSecr=8919380
13	5.978748	10.0.0.2	10.0.0.1	TCP	78	[TCP Dup ACK 11=9] 33228 → 1000 [ACK] Seq=6 Ack=9 Win=29696 Len=0 TSval=8919386 TSecr=8919381 SLE=1
14	6.716803	10.0.0.1	10.0.0.2	TCP	1090	1000 → 33228 [PSH, ACK] Seq=9 Ack=6 Win=29184 Len=1024 TSval=8919565 TSecr=8919386
15	6.716806	10.0.0.1	10.0.0.2	TCP	1514	1000 → 33228 [ACK] Seq=1033 Ack=6 Win=29184 Len=1448 TSval=8919565 TSecr=8919386
16	6.716876	10.0.0.2	10.0.0.1	TCP	66	33228 → 1000 [ACK] Seq=6 Ack=1033 Win=31744 Len=0 TSval=8919570 TSecr=8919565
17	6.716882	10.0.0.2	10.0.0.1	TCP	66	33228 → 1000 [ACK] Seq=6 Ack=2481 Win=34304 Len=0 TSval=8919570 TSecr=8919565
18	6.717614	10.0.0.1	10.0.0.2	TCP	1514	1000 → 33228 [ACK] Seq=2481 Ack=6 Win=29184 Len=1448 TSval=8919565 TSecr=8919386
19	6.717628	10.0.0.2	10.0.0.1	TCP	66	33228 → 1000 [ACK] Seq=6 Ack=3929 Win=37376 Len=0 TSval=8919571 TSecr=8919565
20	6.718804	10.0.0.1	10.0.0.2	TCP	1514	1000 → 33228 [ACK] Seq=3929 Ack=6 Win=29184 Len=1448 TSval=8919565 TSecr=8919386
21	6.718828	10.0.0.2	10.0.0.1	TCP	66	33228 → 1000 [ACK] Seq=6 Ack=5377 Win=39936 Len=0 TSval=8919571 TSecr=8919565
22	6.720019	10.0.0.1	10.0.0.2	TCP	1514	1000 → 33228 [ACK] Seq=5377 Ack=6 Win=29184 Len=1448 TSval=8919565 TSecr=8919386
23	6.720049	10.0.0.2	10.0.0.1	TCP	66	33228 → 1000 [ACK] Seq=6 Ack=6825 Win=43008 Len=0 TSval=8919571 TSecr=8919565
24	6.721292	10.0.0.1	10.0.0.2	TCP	1514	[TCP Previous segment not captured] 1000 → 33228 [ACK] Seq=8273 Ack=6 Win=29184 Len=1448 TSval=8919571 TSecr=8919565
25	6.721329	10.0.0.2	10.0.0.1	TCP	78	[TCP Window Update] 33228 → 1000 [ACK] Seq=6 Ack=6825 Win=46080 Len=0 TSval=8919571 TSecr=8919565 SLE=1
26	6.722588	10.0.0.1	10.0.0.2	TCP	1514	1000 → 33228 [ACK] Seq=9721 Ack=6 Win=29184 Len=1448 TSval=8919566 TSecr=8919386
27	6.722612	10.0.0.2	10.0.0.1	TCP	78	[TCP Window Update] 33228 → 1000 [ACK] Seq=6 Ack=6825 Win=48640 Len=0 TSval=8919572 TSecr=8919565 SLE=1
28	6.723678	10.0.0.1	10.0.0.2	TCP	1514	1000 → 33228 [ACK] Seq=11169 Ack=6 Win=29184 Len=1448 TSval=8919566 TSecr=8919386
29	6.723688	10.0.0.2	10.0.0.1	TCP	78	[TCP Window Update] 33228 → 1000 [ACK] Seq=6 Ack=6825 Win=51712 Len=0 TSval=8919572 TSecr=8919565 SLE=1
30	6.757050	10.0.0.1	10.0.0.2	TCP	1514	[TCP Previous segment not captured] 1000 → 33228 [ACK] Seq=14065 Ack=6 Win=29184 Len=1448 TSval=8919572 TSecr=8919565 SLE=1
31	6.758108	10.0.0.2	10.0.0.1	TCP	86	[TCP Window Update] 33228 → 1000 [ACK] Seq=6 Ack=6825 Win=54784 Len=0 TSval=8919581 TSecr=8919565 SLE=1
32	6.759229	10.0.0.1	10.0.0.2	TCP	1514	1000 → 33228 [ACK] Seq=15513 Ack=6 Win=29184 Len=1448 TSval=8919576 TSecr=8919571
33	6.759275	10.0.0.2	10.0.0.1	TCP	86	[TCP Window Update] 33228 → 1000 [ACK] Seq=6 Ack=6825 Win=57344 Len=0 TSval=8919581 TSecr=8919565 SLE=1
34	6.761779	10.0.0.1	10.0.0.2	TCP	2962	1000 → 33228 [ACK] Seq=16961 Ack=6 Win=29184 Len=2896 TSval=8919577 TSecr=8919571
35	6.761818	10.0.0.2	10.0.0.1	TCP	86	[TCP Window Update] 33228 → 1000 [ACK] Seq=6 Ack=6825 Win=63488 Len=0 TSval=8919582 TSecr=8919565 SLE=1
36	6.763178	10.0.0.1	10.0.0.2	TCP	1514	1000 → 33228 [ACK] Seq=19857 Ack=6 Win=29184 Len=1448 TSval=8919577 TSecr=8919572
37	6.763224	10.0.0.2	10.0.0.1	TCP	86	[TCP Window Update] 33228 → 1000 [ACK] Seq=6 Ack=6825 Win=66048 Len=0 TSval=8919582 TSecr=8919565 SLE=1
38	6.764384	10.0.0.1	10.0.0.2	TCP	1514	[TCP Out-Of-Order] 1000 → 33228 [ACK] Seq=6825 Ack=6 Win=29184 Len=1448 TSval=8919577 TSecr=8919572
39	6.764438	10.0.0.2	10.0.0.1	TCP	78	33228 → 1000 [ACK] Seq=6 Ack=12617 Win=69120 Len=0 TSval=8919582 TSecr=8919577 SLE=14665 SRE=21305
40	6.799678	10.0.0.1	10.0.0.2	TCP	1514	1000 → 33228 [ACK] Seq=21305 Ack=6 Win=29184 Len=1448 TSval=8919586 TSecr=8919581
41	6.799763	10.0.0.2	10.0.0.1	TCP	78	[TCP Window Update] 33228 → 1000 [ACK] Seq=6 Ack=12617 Win=72192 Len=0 TSval=8919591 TSecr=8919577 SLE=14665 SRE=21305
42	6.802733	10.0.0.1	10.0.0.2	TCP	1514	[TCP Out-Of-Order] 1000 → 33228 [ACK] Seq=12617 Ack=6 Win=29184 Len=1448 TSval=8919587 TSecr=8919582
43	6.802735	10.0.0.1	10.0.0.2	TCP	1514	1000 → 33228 [ACK] Seq=22753 Ack=6 Win=29184 Len=1448 TSval=8919587 TSecr=8919582

Figure 20: Part of the packets captured when transmit file with size of 100KB

▶ Frame 8: 71 bytes on wire (568 bits), 71 bytes captured (568 bits)

▶ Ethernet II, Src: 00:00:00:00:00:02 (00:00:00:00:00:02), Dst: 00:00:00:00:00:01 (00:00:00:00:00:01)

▶ Internet Protocol Version 4, Src: 10.0.0.2, Dst: 10.0.0.1

▶ Transmission Control Protocol, Src Port: 33228, Dst Port: 1000, Seq: 1, Ack: 1, Len: 5

▶ Data (5 bytes)

Data: 7465737432

[Length: 5]

```

0000  00 00 00 00 00 01 00 00 00 00 02 08 00 45 00  ....E.
0010  00 39 02 60 40 00 40 06 24 5d 0a 00 00 02 0a 00  9 @ 0 $]
0020  00 01 81 cc 03 e8 26 a6 1a 55 56 44 ef ab 00 18  .&..UVD
0030  00 3a 14 2e 00 00 01 01 08 0a 00 88 19 0e 00 88  .:.....
0040  19 08 74 65 73 74 32  ..Test2

```

No.: 8 · Time: 5.674276 · Source: 10.0.0.2 · Destination: 10.0.0.1 · Protocol: TCP · Length: 71 · Info: 33228 → 1000 [PSH, ACK] Seq=1 Ack=1 Win=29696 Len=5 TSval=8919310 TSecr=8919304

Help Close

Figure 21: Detailed information about packet No.8



Figure 22: Detailed information about packet No.10

No.	Time	Source	Destination	Protocol	Length	Info
1	0.000000	fe80::c800:b9ff:fe..ff02::fb	MDNS	192	Standard query 0x0000 PTR _afpovertcp._tcp.local, "QM" question PTR _ipp._tcp.local, "QM" question ..	
2	0.322544	fe80::fc86:73ff:fe..ff02::fb	MDNS	192	Standard query 0x0000 PTR _afpovertcp._tcp.local, "QM" question PTR _ipp._tcp.local, "QM" question ..	
3	0.363182	fe80::1c1a:bcff:fe..ff02::fb	MDNS	303	Standard query response 0x0000 TXT, cache flush AAAA, cache flush fe80::1c1a:bcff:fe57:d3ac PTR, ca..	
4	0.419983	fe80::28c2:55ff:fe..ff02::fb	MDNS	243	Standard query response 0x0000 PTR _workstation._tcp.local PTR ubuntu [2a:c2:55:73:a3:1a]..workstat..	
5	0.615742	fe80::b089:4fff:fe1..ff02::fb	MDNS	192	Standard query 0x0000 PTR _afpovertcp._tcp.local, "QM" question PTR _ipp._tcp.local, "QM" question ..	
6	0.708538	fe80::c800:b9ff:fe..ff02::fb	MDNS	243	Standard query response 0x0000 PTR _workstation._tcp.local PTR ubuntu [ca:00:b9:14:c5:fb]..workstat..	
7	1.034141	fe80::fc86:73ff:fe..ff02::fb	MDNS	243	Standard query response 0x0000 PTR _workstation._tcp.local PTR ubuntu [fe:86:73:85:cd:ce]..workstat..	
8	1.060555	fe80::28c2:55ff:fe..ff02::fb	MDNS	303	Standard query response 0x0000 TXT, cache flush AAAA, cache flush fe80::28c2:55ff:fe73:a31a PTR, ca..	
9	1.324405	fe80::b089:4fff:fe1..ff02::fb	MDNS	243	Standard query response 0x0000 PTR _workstation._tcp.local PTR ubuntu [b2:89:04:10:d5:39]..workstat..	
10	1.349308	fe80::c800:b9ff:fe..ff02::fb	MDNS	303	Standard query response 0x0000 TXT, cache flush AAAA, cache flush fe80::c800:b9ff:fe14:c5fb PTR, ca..	
11	1.671424	fe80::fc86:73ff:fe..ff02::fb	MDNS	303	Standard query response 0x0000 TXT, cache flush AAAA, cache flush fe80::fc86:73ff:fe85:cdce PTR, ca..	
12	1.963287	fe80::b089:4fff:fe1..ff02::fb	MDNS	303	Standard query response 0x0000 TXT, cache flush AAAA, cache flush fe80::b089:4fff:fe10:d539 PTR, cac..	
13	3.016118	fe80::1c1a:bcff:fe..ff02::fb	MDNS	192	Standard query 0x0000 PTR _afpovertcp._tcp.local, "QM" question PTR _ipp._tcp.local, "QM" question ..	
14	3.714655	fe80::28c2:55ff:fe..ff02::fb	MDNS	192	Standard query 0x0000 PTR _afpovertcp._tcp.local, "QM" question PTR _ipp._tcp.local, "QM" question ..	
15	3.855423	10.0.0.2	10.0.0.1	TCP	74	46382 → 1000 [SYN] Seq=0 Win=29200 Len=0 MSS=1460 SACK_PERM=1 TSval=8922198 TSecr=0 WS=512
16	4.001169	fe80::c800:b9ff:fe..ff02::fb	MDNS	192	Standard query 0x0000 PTR _afpovertcp._tcp.local, "QM" question PTR _ipp._tcp.local, "QM" question ..	
17	4.324479	fe80::fc86:73ff:fe..ff02::fb	MDNS	192	Standard query 0x0000 PTR _afpovertcp._tcp.local, "QM" question PTR _ipp._tcp.local, "QM" question ..	
18	4.617811	fe80::b089:4fff:fe1..ff02::fb	MDNS	192	Standard query 0x0000 PTR _afpovertcp._tcp.local, "QM" question PTR _ipp._tcp.local, "QM" question ..	
19	4.655392	10.0.0.2	10.0.0.1	TCP	74	[TCP Retransmission] 46382 → 1000 [SYN] Seq=0 Win=29200 Len=0 MSS=1460 SACK_PERM=1 TSval=8922198 TSecr=0 WS=512
20	4.897698	10.0.0.1	10.0.0.2	TCP	74	1000 → 46382 [SYN, ACK] Seq=1 Ack=1 Win=29696 Len=0 MSS=1460 SACK_PERM=1 TSval=8922198 TSecr=892219..
21	4.897711	10.0.0.2	10.0.0.1	TCP	66	46382 → 1000 [ACK] Seq=1 Ack=1 Win=29696 Len=0 TSval=8922458 TSecr=8922453
22	4.898018	10.0.0.2	10.0.0.1	TCP	74	46382 → 1000 [PSH, ACK] Seq=1 Ack=1 Win=29696 Len=5 TSval=8922458 TSecr=8922453
23	4.938629	10.0.0.1	10.0.0.2	TCP	66	1000 → 46382 [ACK] Seq=1 Ack=6 Win=29184 Len=0 TSval=8922463 TSecr=8922458
24	4.940548	10.0.0.1	10.0.0.2	TCP	74	1000 → 46382 [PSH, ACK] Seq=1 Ack=6 Win=29184 Len=8 TSval=8922464 TSecr=8922458
25	4.940572	10.0.0.2	10.0.0.1	TCP	66	46382 → 1000 [ACK] Seq=6 Ack=9 Win=29696 Len=0 TSval=8922469 TSecr=8922464
26	5.941133	10.0.0.1	10.0.0.2	TCP	1090	1000 → 46382 [PSH, ACK] Seq=9 Ack=6 Win=29184 Len=1024 TSval=8922714 TSecr=8922469
27	5.941154	10.0.0.2	10.0.0.1	TCP	66	46382 → 1000 [ACK] Seq=6 Ack=1033 Win=31744 Len=0 TSval=8922719 TSecr=8922714
28	5.942647	10.0.0.1	10.0.0.2	TCP	1514	1000 → 46382 [ACK] Seq=1033 Ack=6 Win=29184 Len=1448 TSval=8922714 TSecr=8922469
29	5.942660	10.0.0.2	10.0.0.1	TCP	66	46382 → 1000 [ACK] Seq=6 Ack=2481 Win=34304 Len=0 TSval=8922719 TSecr=8922714
30	5.943111	10.0.0.1	10.0.0.2	TCP	1514	1000 → 46382 [ACK] Seq=2481 Ack=6 Win=29184 Len=1448 TSval=8922714 TSecr=8922469
31	5.943123	10.0.0.2	10.0.0.1	TCP	66	46382 → 1000 [ACK] Seq=6 Ack=3929 Win=37376 Len=0 TSval=8922720 TSecr=8922714
32	5.943851	10.0.0.1	10.0.0.2	TCP	1514	1000 → 46382 [ACK] Seq=3929 Ack=6 Win=29184 Len=1448 TSval=8922715 TSecr=8922469
33	5.943863	10.0.0.2	10.0.0.1	TCP	66	46382 → 1000 [ACK] Seq=6 Ack=5377 Win=39936 Len=0 TSval=8922720 TSecr=8922715
34	5.945150	10.0.0.1	10.0.0.2	TCP	1514	[TCP Previous segment not captured] 1000 → 46382 [ACK] Seq=6825 Ack=6 Win=29184 Len=1448 TSval=8922..
35	5.945162	10.0.0.2	10.0.0.1	TCP	78	[TCP Window Update] 46382 → 1000 [ACK] Seq=6 Ack=5377 Win=43008 Len=0 TSval=8922720 TSecr=8922715 S..
36	5.946211	10.0.0.1	10.0.0.2	TCP	1514	1000 → 46382 [ACK] Seq=8273 Ack=6 Win=29184 Len=1448 TSval=8922715 TSecr=8922469

Figure 23: Part of the packets captured when transmit file with size of 10MB

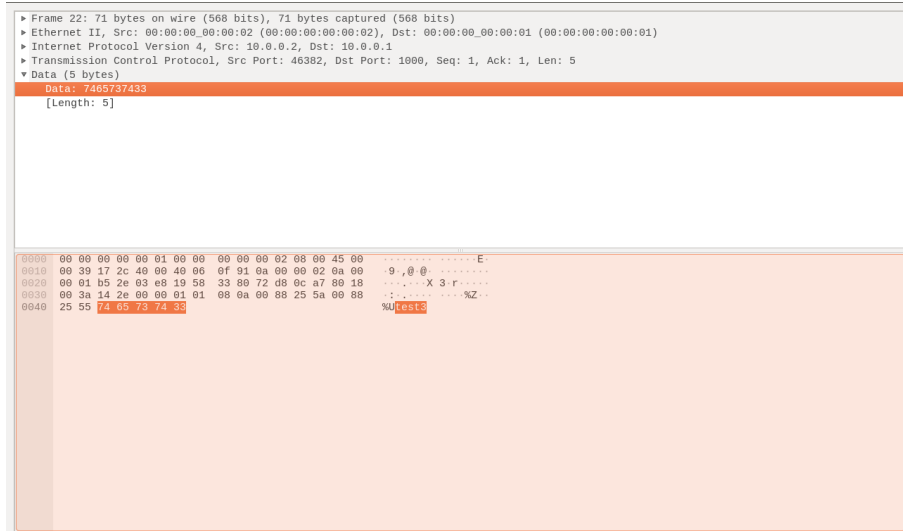


Figure 24: Detailed information about packet No.22



Figure 25: Detailed information about packet No.24

And packet No.24 contains the file length, which is from $h1$ to $h2$. The file length is $0x989680 = 10000000$ bytes = 10MB, which is correct.

And packets with size of more than 1000 bytes contain the data, which is from $h1$ to $h2$. By using `diff`, we can say that the transmission is correct.

3 Part III

3.1 Implement a simple SR

3.1.1 Implementation

The code of `sr.c` and `makeFile` can be seen in **Appendix**.

3.1.2 Test and Results

There're five conditions, in each condition, I use the given file `file3.jpg` as the transferred file, whose size is 4.4MB. I run the sender on `h1` (10.0.0.1) and run the receiver on `h2` (10.0.0.2). The `mitm` program is running on `h3` (10.0.0.3).

A. No reordering, no loss, no error.

By compared the file in receiver side and sender side with `diff`, the result is correct. The first few lines of log on sender is as following:

```
SYN 234756642 0 0
ACK 234756642 0 0
DATA 0 1456 3220623233
DATA 1 1456 2244073028
DATA 2 1456 1460592462
DATA 3 1456 3415882852
DATA 4 1456 3456527178
DATA 5 1456 1573440664
DATA 6 1456 3472368927
DATA 7 1456 3178870208
DATA 8 1456 279932998
DATA 9 1456 3426093530
ACK 1 0 0
```

Then, the window won't slide when `DATA` packets is sent, and the window will slide one packet when `ACK 1 0 0` comes.

The first few lines of log on receiver is as following:

```
SYN 234756642 0 0
ACK 234756642 0 0
DATA 0 1456 3220623233
ACK 1 0 0
DATA 1 1456 2244073028
ACK 2 0 0
DATA 2 1456 1460592462
ACK 3 0 0
DATA 3 1456 3415882852
```

Then, the window won't slide when `ACK` packets is sent. And in this case, since each `Data` packets comes with desired seq number, the window will slide for one packet for each `DATA` packets.

B. 10 % loss, no reordering, no error

By compared the file in receiver side and sender side with `diff`, the result is correct. As for the sender side, the evidence of retransmission can be seen in the following lines:

```
DATA 47 1456 2926046941
ACK 38 0 0
ACK 38 0 0
ACK 38 0 0
ACK 38 0 0
ACK 38 0 0
ACK 38 0 0
ACK 38 0 0
DATA 38 1456 2564180116
DATA 39 1456 214759100
DATA 40 1456 1477671586
DATA 41 1456 4058279732
DATA 42 1456 3378489669
DATA 43 1456 10853198
DATA 44 1456 1483690205
DATA 45 1456 2820064048
DATA 46 1456 720727197
DATA 47 1456 2926046941
ACK 42 0 0
```

It can be seen that at least DATA 38 is lost. And after retransmission, the packets is ACKed with next seq number 42.

As for the receiver side, the evidence of packet loss can be seen in the following lines:

```
DATA 37 1456 1168804078
ACK 38 0 0
DATA 39 1456 214759100
ACK 38 0 0
DATA 40 1456 1477671586
ACK 38 0 0
DATA 41 1456 4058279732
ACK 38 0 0
DATA 43 1456 10853198
ACK 38 0 0
DATA 44 1456 1483690205
ACK 38 0 0
DATA 45 1456 2820064048
ACK 38 0 0
DATA 46 1456 720727197
```


ACK 38 0 0
DATA 47 1456 2926046941
ACK 38 0 0
DATA 38 1456 2564180116
ACK 42 0 0

Then, it can be seen that DATA 38 is the gap.

Thus, DATA 38 is lost.

C. 10 % error, no reordering, no loss

By compared the file in receiver side and sender side with `diff`, the result is correct. As for the sender side, the evidence of retransmission can be seen in the following lines:

DATA 10 1456 2805442653
ACK 2 0 0
DATA 11 1456 443311922
ACK 3 0 0
DATA 12 1456 4027845571
ACK 4 0 0
DATA 13 1456 1330974691
ACK 5 0 0
DATA 14 1456 2109819272
ACK 6 0 0
DATA 15 1456 965284191
ACK 7 0 0
DATA 16 1456 3311776188
ACK 8 0 0
DATA 17 1456 1539359076
ACK 9 0 0
DATA 18 1456 2360034080
ACK 10 0 0
DATA 19 1456 725513889
ACK 10 0 0
ACK 10 0 0
ACK 10 0 0
ACK 10 0 0
ACK 10 0 0
ACK 10 0 0
ACK 10 0 0
ACK 10 0 0
ACK 10 0 0

```
ACK 10 0 0
DATA 10 1456 2805442653
DATA 11 1456 443311922
DATA 12 1456 4027845571
DATA 13 1456 1330974691
DATA 14 1456 2109819272
DATA 15 1456 965284191
DATA 16 1456 3311776188
DATA 17 1456 1539359076
DATA 18 1456 2360034080
DATA 19 1456 725513889
ACK 20 0 0
```

It can be seen that at least DATA 10 is corrupted since it has been transmitted but not ACKed. And after retransmission, the packets is ACKed with next seq number 20.

As for the receiver side, the evidence of packet loss can be seen in the following lines:

```
ACK 10 0 0
DATA 10 1456 2805442653
DATA 11 1456 443311922
ACK 10 0 0
```

Then, it can be seen that DATA 10 is received but not ACKed due to corruption.

Thus, DATA 10 is corrupted

D. Reordering, no error, no loss

By compared the file in receiver side and sender side with `diff`, the result is correct. The first few lines for log on sender is:

```
SYN 1194581707 0 0
ACK 1194581707 0 0
DATA 0 1456 3220623233
DATA 1 1456 2244073028
DATA 2 1456 1460592462
DATA 3 1456 3415882852
DATA 4 1456 3456527178
DATA 5 1456 1573440664
DATA 6 1456 3472368927
DATA 7 1456 3178870208
DATA 8 1456 279932998
DATA 9 1456 3426093530
ACK 0 0 0
```

```
ACK 0 0 0
ACK 0 0 0
ACK 0 0 0
ACK 0 0 0
ACK 0 0 0
ACK 0 0 0
ACK 8 0 0
DATA 10 1456 2805442653
ACK 10 0 0
```

Since ACK 10 0 0 is received, we can say that the first 10 packets are delivered. And clearly, shuffle happens and make the first data packet to receiver is not DATA 0.

The first few lines for log on receiver is:

```
SYN 1194581707 0 0
ACK 1194581707 0 0
DATA 4 1456 3456527178
ACK 0 0 0
DATA 3 1456 3415882852
ACK 0 0 0
DATA 5 1456 1573440664
ACK 0 0 0
DATA 2 1456 1460592462
ACK 0 0 0
DATA 6 1456 3472368927
ACK 0 0 0
DATA 1 1456 2244073028
ACK 0 0 0
DATA 7 1456 3178870208
ACK 0 0 0
DATA 0 1456 3220623233
ACK 8 0 0
DATA 8 1456 279932998
ACK 9 0 0
DATA 9 1456 3426093530
ACK 10 0 0
```

As we can see, the order of data is completely re-ordered. So, how does receiver and sender make it correct. For receiver, it has a buffer. And the order within a buffer is correct, which ensures the correctness. As for a sender, A later ACK seq number may be shuffled to a recent timing, but since the

ACK means the packets before has been received, so it won't brought trouble.

E. Reordering, 5% loss, 5% error

By compared the file in receiver side and sender side with `diff`, the result is correct.

3.2 Make sender more efficient – leverage duplicate ACK on sender side

3.2.1 Implementation

The code of `sr.c` and `makeFile` can be seen in **Appendix**.

3.2.2 Tests and Results

I use the given file `file3.jpg` as the transferred file, whose size is 4.4MB. I run the sender on `h1` (10.0.0.1) and run the receiver on `h2` (10.0.0.2). The `mitm` program is running on `h3` (10.0.0.3).

A. 10% loss, no reordering, no error.

The following lines indicates that receive three duplicate ACKs and resend DATA 1.

```
ACK 1 0 0
DATA 10 1456 2805442653
ACK 1 0 0
ACK 1 0 0
DATA 1 1456 2244073028
```

B. Efficiency.

The transferring time of sender in part 3.1 is 1min51.650s and the transferring time of sender in part 3.2 is 1min35.900s Duplicate ACK make the sender more efficient.

3.3 Make sender more efficient – send NACK on receiver side

3.3.1 Implementation

The code of `sr.c` and `makeFile` can be seen in **Appendix**.

3.3.2 Tests and Results

I use the given file `file3.jpg` as the transferred file, whose size is 4.4MB. I run the sender on `h1` (10.0.0.1) and run the receiver on `h2` (10.0.0.2). The `mitm` program is running on `h3` (10.0.0.3).

A. 10% loss, no reordering, no error.

The following lines indicates that receive three duplicate ACKs and resend DATA 1.

```
ACK 1 0 0
DATA 10 1456 2805442653
ACK 1 0 0
ACK 1 0 0
```

DATA 1 1456 2244073028

B. Efficiency.

The transferring time of sender in part 3.1 is 1min51.650s, and the transferring time of sender in part 3.3 is 1min30.650s. NACK make the sender more efficient.

3.4 Throughput, delay and window size

Throughput under different condition can be seen in Table 1 It can be seen that the throughput has

	delay = 0.01ms	delay = 10ms	delay = 50ms	delay = 100ms
window size = 10	2.722s	7.157s	31.871s	1min1.892s
window size = 50	0.758s	1.695s	6.934s	

Table 1: Throughput at different conditions

positive linear relationship with the inverse of the window size. And when the delay is not very small. the throughput seems has positive linear relationship with delay. When the delay is too small, other operation like file IO will influence the throughput greatly.

4 Appendix

4.1 Part II

4.1.1 makeFile

```
1  # use c++ 11 standard here (c++17 would be fine?)
2  CC=g++ -g -Wall -std=c++11
3
4  # List of source files for iPerfer
5  FS_SOURCES=ftrans.cc
6
7  # Generate the names of the iPerfer's object files
8  FS_OBJS=${FS_SOURCES%.cc=.o}
9
10 all: ftrans
11
12 ftrans: ${FS_OBJS}
13     ${CC} -o $@ $^
14
15 # Generic rules for compiling a source file to an object file
16 %.o: %.cpp
17     ${CC} -c $<
18 %.o: %.cc
19     ${CC} -c $<
20
21 clean:
22     rm -f ${FS_OBJS} ftrans
```

4.1.2 ftrans.c

```
1  # use c++ 11 standard here (c++17 would be fine?)
2  CC=g++ -g -Wall -std=c++11
3
4  # List of source files for iPerfer
5  FS_SOURCES=ftrans.cc
```

```

6
7  # Generate the names of the iPerfer's object files
8  FS_OBJS=${FS_SOURCES%.cc=.o}
9
10 all: ftrans
11
12 ftrans: ${FS_OBJS}
13     ${CC} -o $@ $^
14
15 # Generic rules for compiling a source file to an object file
16 %.o: %.cpp
17     ${CC} -c $<
18 %.o: %.cc
19     ${CC} -c $<
20
21 clean:
22     rm -f ${FS_OBJS} ftrans

```

4.2 Part III

4.2.1 makeFile

```

1  CC = gcc
2
3  CFLAGS = -g -Wall -std=c99 -D_GNU_SOURCE -pthread
4
5  # the build target executable:
6  SR = sr
7  MITM = mitm
8
9  all: $(SR) $(MITM)
10
11 # assume the C code for TARGET is TARGET.c
12 $(SR): $(SR).c
13     $(CC) $(CFLAGS) -o $(SR) $(SR).c
14
15 $(MITM): $(MITM).c
16     $(CC) $(CFLAGS) -o $(MITM) $(MITM).c
17
18 clean:
19     $(RM) $(SR) $(MITM)

```

4.2.2 sr.c in 3.1

```

1  #include "SRHeader.h"
2  #include "crc32.h"
3  #include <arpa/inet.h>
4  #include <netdb.h>
5  #include <stdint.h>
6  #include <stdio.h>
7  #include <stdlib.h>
8  #include <string.h>
9  #include <sys/socket.h>
10 #include <sys/stat.h> // struct stat
11 #include <sys/time.h>
12 #include <time.h>
13 #include <unistd.h>
14 // The size of packet would is len(header) + len(payload) = 4*4 + 1456 = 1472
15 #define MAX_MESSAGE_SIZE 1472
16
17 int logInfo(struct SRHeader header, FILE *fp) {
18     if (header.flag == SYN) {

```

```

19     fprintf(fp, "%s %u %u %u\n", "SYN", header.seq, header.len, header.crc);
20     printf("%s %u %u %u\n", "SYN", header.seq, header.len, header.crc);
21 } else if (header.flag == FIN) {
22     fprintf(fp, "%s %u %u %u\n", "FIN", header.seq, header.len, header.crc);
23     printf("%s %u %u %u\n", "FIN", header.seq, header.len, header.crc);
24 } else if (header.flag == DATA) {
25     fprintf(fp, "%s %u %u %u\n", "DATA", header.seq, header.len, header.crc);
26     printf("%s %u %u %u\n", "DATA", header.seq, header.len, header.crc);
27 } else if (header.flag == ACK) {
28     fprintf(fp, "%s %u %u %u\n", "ACK", header.seq, header.len, header.crc);
29     printf("%s %u %u %u\n", "ACK", header.seq, header.len, header.crc);
30 } else if (header.flag == NACK) {
31     fprintf(fp, "%s %u %u %u\n", "NACK", header.seq, header.len, header.crc);
32     printf("%s %u %u %u\n", "NACK", header.seq, header.len, header.crc);
33 }
34 return 0;
35 }
36
37 int runSender(char *recvIP, int recvPort, int senderPort,
38             unsigned int windowSize, char *fileToSend, char *logFile) {
39     FILE *logfp = fopen(logFile, "w+");
40     // create UDP socket.
41     int sockfd = socket(AF_INET, SOCK_DGRAM, IPPROTO_UDP); // use IPPROTO_UDP
42     if (sockfd < 0) {
43         perror("socket creation failed");
44         exit(1);
45     }
46
47     // below is the same as TCP socket.
48     struct sockaddr_in myAddr;
49     memset(&myAddr, 0, sizeof(myAddr));
50     myAddr.sin_family = AF_INET;
51     myAddr.sin_addr.s_addr = INADDR_ANY;
52     // bind to a specific port
53     myAddr.sin_port = htons(senderPort);
54     bind(sockfd, (struct sockaddr *)&myAddr, sizeof(myAddr));
55
56     struct sockaddr_in recvAddr;
57     struct hostent *host = gethostbyname(recvIP);
58     memcpy(&(recvAddr.sin_addr), host->h_addr, host->h_length);
59     recvAddr.sin_family = AF_INET;
60     recvAddr.sin_port = htons(recvPort);
61
62     // 1. The num of packets to be sent
63     // 1.1 Read file length
64     struct stat statBuffer;
65     size_t fileLen = 0;
66     size_t numSend = 0;
67     size_t divFlag = 0;
68     // Read the file length
69     if (stat(fileToSend, &statBuffer) < 0) {
70         // Handling reding file error
71         perror("Error: Can not read file\n");
72         exit(-1);
73     }
74     fileLen = statBuffer.st_size;
75     // 1.2 Calculate num of packets to be sent
76     divFlag = fileLen % MAX_PAYLOAD_SIZE;
77     numSend = divFlag > 0 ? (1 + (fileLen / MAX_PAYLOAD_SIZE))
78                 : fileLen / MAX_PAYLOAD_SIZE;
79     printf("File size is %lu, need %lu packets to send\n", fileLen, numSend);
80
81     // 2. send

```



```

82     unsigned int randSeq = 0;
83     char packet[MAX_MESSAGE_SIZE];
84     char msg[MAX_MESSAGE_SIZE];
85     struct SRHeader header;
86     memset(packet, 0, sizeof(packet));
87     // 2.1 send SYN, has random seq number
88     srand((unsigned)time(NULL));
89     randSeq = rand();
90     header.flag = SYN;
91     header.seq = randSeq;
92     header.len = 0;
93     header.crc = crc32(NULL, 0); // without payload;
94     memcpy(packet, (char *)&header, sizeof(header));
95     sendto(sockfd, (void *)packet, sizeof(header), 0,
96            (const struct sockaddr *)&recvAddr, sizeof(struct sockaddr_in));
97     logInfo(header, logfp);
98     printf("Send SYN to set up the connection\n");
99     // 2.2 wait for ACK and SYN packet
100    socklen_t sLen = sizeof(struct sockaddr_in);
101    struct sockaddr_in siOther;
102    struct SRHeader recvHeader;
103    struct timeval start, stop;
104    gettimeofday(&start, NULL);
105    gettimeofday(&stop, NULL);
106
107    // receive ack
108    while (1) {
109        if (recvfrom(sockfd, msg, MAX_MESSAGE_SIZE, MSG_DONTWAIT,
110                    (struct sockaddr *)&siOther, &sLen) > 0) {
111            memcpy(&recvHeader, msg, sizeof(recvHeader));
112            logInfo(recvHeader, logfp);
113            if (recvHeader.flag == ACK) {
114                printf("Connection set up \n");
115                break;
116            }
117        } else {
118            // resend SYN
119            sendto(sockfd, (void *)packet, sizeof(header), 0,
120                   (const struct sockaddr *)&recvAddr, sizeof(struct sockaddr_in));
121            logInfo(header, logfp);
122        }
123        // time out for ACK
124        gettimeofday(&stop, NULL);
125        if ((double)(stop.tv_sec - start.tv_sec) * 1000 +
126            (double)(stop.tv_usec - start.tv_usec) / 1000 >
127            5000) {
128            perror("Error: do not receive ACK for SYN within 5 s. Give up\n");
129            fclose(logfp);
130            exit(4);
131        }
132    }
133    // 2.3 Send data
134    char *buf = malloc(windowSize * MAX_MESSAGE_SIZE);
135    memset(buf, 0, windowSize * MAX_MESSAGE_SIZE);
136    FILE *fp = fopen(fileToSend, "r");
137    // open file
138    if (!fp) {
139        // Handling file open error
140        perror("Error: Can not read file\n");
141        free(buf);
142        fclose(logfp);
143        fclose(fp);
144        exit(-1);

```

```

145     }
146
147     unsigned int currentSeq = 0; // send 0 for initial
148     unsigned int bufNum = 0;
149     int endFlag = 1;
150     size_t payloadLen = MAX_PAYLOAD_SIZE;
151     gettimeofday(&start, NULL);
152     while (endFlag) {
153         // check if timeout
154         gettimeofday(&stop, NULL);
155         // if timeout, resend all buffer
156         if ((double)(stop.tv_sec - start.tv_sec) * 1000 +
157             (double)(stop.tv_usec - start.tv_usec) / 1000 >
158             400) {
159             printf(
160                 "No ack within 400ms, resend all packets in buffer. Start from %u\n",
161                 currentSeq);
162             // resend
163             for (unsigned int ii = 0; ii < bufNum; ii++) {
164                 payloadLen = ((currentSeq + ii == numSend - 1) && (divFlag > 0))
165                     ? divFlag
166                     : MAX_PAYLOAD_SIZE;
167                 sendto(sockfd, (void *)(buf + ii * sizeof(packet)),
168                     payloadLen + sizeof(header), 0,
169                     (const struct sockaddr *)&recvAddr, sizeof(struct sockaddr_in));
170                 struct SRHeader tempHeader;
171                 memcpy(&tempHeader, buf + ii * sizeof(packet), sizeof(header));
172                 logInfo(tempHeader, logfp);
173             }
174             // reset time
175             gettimeofday(&start, NULL);
176         }
177         // send and buffer
178
179         if ((bufNum < windowSize) && ((bufNum + currentSeq) < numSend)) {
180             payloadLen = ((currentSeq + bufNum == numSend - 1) && (divFlag > 0))
181                 ? divFlag
182                 : MAX_PAYLOAD_SIZE;
183             // form a packet
184             fread(packet + sizeof(header), payloadLen, 1, fp);
185             header.flag = DATA;
186             header.seq = currentSeq + bufNum;
187             header.len = payloadLen;
188             header.crc = crc32(packet + sizeof(header), payloadLen);
189             memcpy(packet, (char *)&header, sizeof(header));
190             // send
191             sendto(sockfd, (void *)packet, payloadLen + sizeof(header), 0,
192                 (const struct sockaddr *)&recvAddr, sizeof(struct sockaddr_in));
193             logInfo(header, logfp);
194             // buffer the packet
195             memcpy(buf + bufNum * MAX_MESSAGE_SIZE, packet, sizeof(packet));
196             bufNum++;
197         }
198         // receive ACK
199         if (recvfrom(sockfd, msg, MAX_MESSAGE_SIZE, MSG_DONTWAIT,
200             (struct sockaddr *)&siOther, &sLen) != -1) {
201             memcpy(&recvHeader, msg, sizeof(recvHeader));
202             logInfo(recvHeader, logfp);
203             // Advance window
204             if (recvHeader.flag == ACK) {
205                 if (recvHeader.seq > currentSeq) {
206                     // move buffer
207                     if (recvHeader.seq > currentSeq + bufNum) {

```

```

208         // clear buf
209         memset(buf, 0, windowSize * MAX_MESSAGE_SIZE);
210         bufNum = 0;
211     } else {
212         char *temp = malloc(windowSize * MAX_PAYLOAD_SIZE);
213         memcpy(temp, buf + (recvHeader.seq - currentSeq) * MAX_MESSAGE_SIZE,
214                (bufNum + currentSeq - recvHeader.seq) * MAX_MESSAGE_SIZE);
215         memset(buf, 0, windowSize * MAX_MESSAGE_SIZE);
216         memcpy(buf, temp, windowSize * MAX_MESSAGE_SIZE);
217         free(temp);
218         bufNum = bufNum + currentSeq - recvHeader.seq;
219     }
220     // advance
221     currentSeq = recvHeader.seq;
222     if (currentSeq > (numSend - 1))
223         endFlag = 0;
224     // reset time
225     gettimeofday(&start, NULL);
226 }
227 }
228 }
229 }
230 // Send FIN
231 header.flag = FIN;
232 header.seq = randSeq;
233 header.len = 0;
234 header.crc = crc32(NULL, 0); // without payload;
235 memcpy(packet, (char *)&header, sizeof(header));
236 sendto(sockfd, (void *)packet, sizeof(header), 0,
237         (const struct sockaddr *)&recvAddr, sizeof(struct sockaddr_in));
238 logInfo(header, logfp);
239 printf("Transmission finished: Send FIN");
240 // receive ack for FIN
241 while (1) {
242     if (recvfrom(sockfd, msg, MAX_MESSAGE_SIZE, MSG_DONTWAIT,
243                 (struct sockaddr *)&siOther, &sLen) != -1) {
244         memcpy(&recvHeader, msg, sizeof(recvHeader));
245         logInfo(recvHeader, logfp);
246         if (recvHeader.flag == ACK) {
247             printf("Receive ACK for FIN, done\n");
248             break;
249         }
250     } else {
251         // resend FIN
252         sendto(sockfd, (void *)packet, sizeof(header), 0,
253               (const struct sockaddr *)&recvAddr, sizeof(struct sockaddr_in));
254         logInfo(header, logfp);
255     }
256     // time out for ACK
257     gettimeofday(&stop, NULL);
258     if ((double)(stop.tv_sec - start.tv_sec) * 1000 +
259         (double)(stop.tv_usec - start.tv_usec) / 1000 >
260         5000) {
261         perror("Error: do not receive ACK for FIN within 5s. Give up\n");
262         fclose(fp);
263         fclose(logfp);
264         free(buf);
265         exit(4);
266     }
267 }
268 fclose(fp);
269 fclose(logfp);
270 free(buf);

```

```

271     return 0;
272 }
273
274 int runReceiver(int port, unsigned int windowSize, char *recvDir,
275               char *logFile) {
276     FILE *logfp = fopen(logFile, "w+");
277     // Set up connection
278     struct sockaddr_in siMe;
279     int sockfd = socket(AF_INET, SOCK_DGRAM, IPPROTO_UDP);
280     if (sockfd == -1) {
281         fputs("socket creation failed!", stderr);
282         exit(2); // error, exit with 2
283     }
284     memset((char *)&siMe, 0, sizeof(struct sockaddr_in));
285     siMe.sin_family = AF_INET;
286     siMe.sin_port = htons(port);
287     siMe.sin_addr.s_addr = htonl(INADDR_ANY);
288     bind(sockfd, (struct sockaddr *)&siMe, sizeof(struct sockaddr_in));
289
290     socklen_t sLen = sizeof(struct sockaddr_in);
291     struct sockaddr_in siOther;
292     struct timeval start, stop;
293     int connectCnt = 0;
294     while (1) {
295         // 1. Ready for receiving SYN
296         char msg[MAX_MESSAGE_SIZE];
297         memset(msg, 0, MAX_MESSAGE_SIZE);
298         char ACKMsg[MAX_MESSAGE_SIZE];
299         memset(ACKMsg, 0, sizeof(ACKMsg));
300         unsigned SYNFINSeq = 0;
301         struct SRHeader recvHeader;
302         struct SRHeader ACKHeader;
303
304         gettimeofday(&start, NULL);
305         gettimeofday(&stop, NULL);
306         while (1) {
307             if (recvfrom(sockfd, msg, MAX_MESSAGE_SIZE, MSG_DONTWAIT,
308                         (struct sockaddr *)&siOther, &sLen) > 0) {
309                 memcpy(&recvHeader, msg, sizeof(recvHeader));
310                 logInfo(recvHeader, logfp);
311                 if (recvHeader.flag == SYN) {
312                     SYNFINSeq = recvHeader.seq;
313                     // ACK for SYN
314                     ACKHeader.flag = ACK;
315                     ACKHeader.seq = SYNFINSeq;
316                     ACKHeader.len = 0;
317                     ACKHeader.crc = crc32(NULL, 0);
318                     memcpy(ACKMsg, (char *)&ACKHeader, sizeof(ACKHeader));
319                     sendto(sockfd, (void *)ACKMsg, sizeof(ACKHeader), 0,
320                           (const struct sockaddr *)&siOther, sizeof(struct sockaddr_in));
321                     logInfo(ACKHeader, logfp);
322                     break;
323                 }
324             }
325             gettimeofday(&stop, NULL);
326             if ((double)(stop.tv_sec - start.tv_sec) * 1000 +
327                 (double)(stop.tv_usec - start.tv_usec) / 1000 >
328                 5000) {
329                 printf("I do not receive SYN within 5 s. Give up\n");
330                 return 0;
331             }
332         }
333         // 2. Recive data

```

```

334
335 // 2.1 open file
336 // create file name
337 char fileName[30];
338 char name[13];
339 sprintf(name, "file_%d.txt", connectCnt);
340 strcpy(fileName, recvDir);
341 strcat(fileName, name);
342 FILE *fp = fopen(fileName, "w");
343 unsigned int currentACKSeq = 0;
344 char *buf = malloc(windowSize * (MAX_PAYLOAD_SIZE + 8)); // valid. len
345 memset(buf, 0, windowSize * (MAX_PAYLOAD_SIZE + 8));
346 int FINflag = 1; // 0 for receive FIN flag
347
348 // 2.2 receive data
349 printf("Start to receive file\n");
350 while (FINflag) {
351     // receive data
352     if (recvfrom(sockfd, msg, MAX_MESSAGE_SIZE, MSG_DONTWAIT,
353         (struct sockaddr *)&siOther, &sLen) != -1) {
354         memcpy(&recvHeader, msg, sizeof(recvHeader));
355         logInfo(recvHeader, logfp);
356         // If it's data
357         if (recvHeader.flag == DATA) {
358             if (recvHeader.crc ==
359                 (crc32(msg + sizeof(recvHeader), recvHeader.len))) {
360                 // check the condition for window operation
361                 if (recvHeader.seq >= currentACKSeq + windowSize) {
362                     continue;
363                 } else if (recvHeader.seq < currentACKSeq) {
364                     // send ACK currentACKSeq
365                     ACKHeader.flag = ACK;
366                     ACKHeader.seq = currentACKSeq;
367                     ACKHeader.len = 0;
368                     ACKHeader.crc = crc32(NULL, 0);
369                     memcpy(ACKMsg, (char *)&ACKHeader, sizeof(ACKHeader));
370                     sendto(sockfd, (void *)ACKMsg, sizeof(ACKHeader), 0,
371                         (const struct sockaddr *)&siOther,
372                         sizeof(struct sockaddr_in));
373                     logInfo(ACKHeader, logfp);
374                 } else if ((recvHeader.seq > currentACKSeq) &&
375                     (recvHeader.seq < currentACKSeq + windowSize)) {
376                     // buffer it
377                     unsigned int temp = 1;
378                     memcpy(buf + (recvHeader.seq - currentACKSeq) *
379                         (MAX_PAYLOAD_SIZE + 8),
380                         (char *)&temp, sizeof(temp));
381                     memcpy(buf +
382                         (recvHeader.seq - currentACKSeq) *
383                         (MAX_PAYLOAD_SIZE + 8) +
384                         4,
385                         (char *)&recvHeader.len, sizeof(temp));
386                     memcpy(buf +
387                         + (recvHeader.seq - currentACKSeq) *
388                         (MAX_PAYLOAD_SIZE + 8) +
389                         8,
390                         msg + sizeof(recvHeader), recvHeader.len);
391                     // send current ACK
392                     ACKHeader.flag = ACK;
393                     ACKHeader.seq = currentACKSeq;
394                     ACKHeader.len = 0;
395                     ACKHeader.crc = crc32(NULL, 0);
396                     memcpy(ACKMsg, (char *)&ACKHeader, sizeof(ACKHeader));

```

```

397         sendto(sockfd, (void *)ACKMsg, sizeof(ACKHeader), 0,
398                     (const struct sockaddr *)&siOther,
399                     sizeof(struct sockaddr_in));
400         logInfo(ACKHeader, logfp);
401     } else if (recvHeader.seq == currentACKSeq) {
402         // calculate the highest available packet
403         unsigned int numWrite = 1;
404         for (numWrite = 1; numWrite < windowSize; numWrite++) {
405             unsigned int temp;
406             memcpy(&temp, buf + numWrite * (8 + MAX_PAYLOAD_SIZE),
407                 sizeof(temp));
408             if (temp == 0)
409                 break;
410         }
411         fwrite(msg + sizeof(recvHeader), 1, recvHeader.len, fp);
412         for (unsigned int ii = 1; ii < numWrite; ii++) {
413             unsigned int tempLen;
414             memcpy(&tempLen, buf + ii * (8 + MAX_PAYLOAD_SIZE) + 4,
415                 sizeof(tempLen));
416             fwrite(buf + ii * (8 + MAX_PAYLOAD_SIZE) + 8, 1, tempLen, fp);
417         }
418         // clear buff
419         memset(buf, 0, windowSize * (MAX_PAYLOAD_SIZE + 4));
420         // ACK
421         currentACKSeq = currentACKSeq + numWrite;
422         // send current ACK
423         ACKHeader.flag = ACK;
424         ACKHeader.seq = currentACKSeq;
425         ACKHeader.len = 0;
426         ACKHeader.crc = crc32(NULL, 0);
427         memcpy(ACKMsg, (char *)&ACKHeader, sizeof(ACKHeader));
428         sendto(sockfd, (void *)ACKMsg, sizeof(ACKHeader), 0,
429             (const struct sockaddr *)&siOther,
430             sizeof(struct sockaddr_in));
431         logInfo(ACKHeader, logfp);
432     }
433 }
434 } else if (recvHeader.flag == FIN) {
435     FINflag = 0;
436     // send ACK
437     ACKHeader.flag = ACK;
438     ACKHeader.seq = SYNFINSeq;
439     ACKHeader.len = 0;
440     ACKHeader.crc = crc32(NULL, 0);
441     memcpy(ACKMsg, (char *)&ACKHeader, sizeof(ACKHeader));
442     sendto(sockfd, (void *)ACKMsg, sizeof(ACKHeader), 0,
443         (const struct sockaddr *)&siOther, sizeof(struct sockaddr_in));
444     logInfo(ACKHeader, logfp);
445     // end
446     fclose(fp);
447     free(buf);
448     connectCnt++;
449     fclose(logfp);
450 }
451 }
452 }
453 }
454 fclose(logfp);
455 return 0;
456 }
457
458 int main(int argc, char **argv) {
459     // Sender mode

```

```

460     if (strcmp(argv[1], "-s") == 0) {
461         // Error for number of arguments
462         if (argc != 8) {
463             perror(
464                 "./sr -s <receiver's IP> <receiver's port> <sender's port> <window "
465                 "size> <file to send> <log file>");
466             return 1; // error: exit with 1
467         }
468         char *recvIP = argv[2];
469         int recvPort = atoi(argv[3]);
470         int senderPort = atoi(argv[4]);
471         unsigned int windowSize = atoi(argv[5]);
472         char *fileToSend = argv[6];
473         char *logFile = argv[7];
474         runSender(recvIP, recvPort, senderPort, windowSize, fileToSend, logFile);
475         return 0;
476     }
477
478     // Receiver mode
479     else if (strcmp(argv[1], "-r") == 0) {
480         // Error for number of arguments
481         if (argc != 6) {
482             perror("./sr -r <port> <window size> <recv dir> <log file>");
483             return 1; // error: exit with 1
484         }
485         // int runReceiver(int port, unsigned int windowSize, char *recvDir,
486         //                 char *logFile)
487         int port = atoi(argv[2]);
488         unsigned int windowSize = atoi(argv[3]);
489         char *recvDir = argv[4];
490         char *logFile = argv[5];
491         runReceiver(port, windowSize, recvDir, logFile);
492         return 0;
493     } else {
494         perror("./sr -s <receiver's IP> <receiver's port> <sender's port> <window "
495             "size> <file to send> <log file>\n ./sr -r <port> <window size> "
496             "<recv dir> <log file>\n");
497         return 1;
498     }
499     return 0;
500 }

```

4.2.3 sr.c in 3.2

```

1  #include "SRHeader.h"
2  #include "crc32.h"
3  #include <arpa/inet.h>
4  #include <netdb.h>
5  #include <stdint.h>
6  #include <stdio.h>
7  #include <stdlib.h>
8  #include <string.h>
9  #include <sys/socket.h>
10 #include <sys/stat.h> // struct stat
11 #include <sys/time.h>
12 #include <time.h>
13 #include <unistd.h>
14 // The size of packet would is len(header) + len(payload) = 4*4 + 1456 = 1472
15 #define MAX_MESSAGE_SIZE 1472
16
17 int logInfo(struct SRHeader header, FILE *fp) {
18     if (header.flag == SYN) {
19         fprintf(fp, "%s %u %u %u\n", "SYN", header.seq, header.len, header.crc);

```



```

20     printf("%s %u %u %u\n", "SYN", header.seq, header.len, header.crc);
21 } else if (header.flag == FIN) {
22     fprintf(fp, "%s %u %u %u\n", "FIN", header.seq, header.len, header.crc);
23     printf("%s %u %u %u\n", "FIN", header.seq, header.len, header.crc);
24 } else if (header.flag == DATA) {
25     fprintf(fp, "%s %u %u %u\n", "DATA", header.seq, header.len, header.crc);
26     printf("%s %u %u %u\n", "DATA", header.seq, header.len, header.crc);
27 } else if (header.flag == ACK) {
28     fprintf(fp, "%s %u %u %u\n", "ACK", header.seq, header.len, header.crc);
29     printf("%s %u %u %u\n", "ACK", header.seq, header.len, header.crc);
30 } else if (header.flag == NACK) {
31     fprintf(fp, "%s %u %u %u\n", "NACK", header.seq, header.len, header.crc);
32     printf("%s %u %u %u\n", "NACK", header.seq, header.len, header.crc);
33 }
34 return 0;
35 }
36
37 int runSender(char *recvIP, int recvPort, int senderPort,
38             unsigned int windowSize, char *fileToSend, char *logFile) {
39     FILE *logfp = fopen(logFile, "w+");
40     // create UDP socket.
41     int sockfd = socket(AF_INET, SOCK_DGRAM, IPPROTO_UDP); // use IPPROTO_UDP
42     if (sockfd < 0) {
43         perror("socket creation failed");
44         exit(1);
45     }
46
47     // below is the same as TCP socket.
48     struct sockaddr_in myAddr;
49     memset(&myAddr, 0, sizeof(myAddr));
50     myAddr.sin_family = AF_INET;
51     myAddr.sin_addr.s_addr = INADDR_ANY;
52     // bind to a specific port
53     myAddr.sin_port = htons(senderPort);
54     bind(sockfd, (struct sockaddr *)&myAddr, sizeof(myAddr));
55
56     struct sockaddr_in recvAddr;
57     struct hostent *host = gethostbyname(recvIP);
58     memcpy(&(recvAddr.sin_addr), host->h_addr, host->h_length);
59     recvAddr.sin_family = AF_INET;
60     recvAddr.sin_port = htons(recvPort);
61
62     // 1. The num of packets to be sent
63     // 1.1 Read file length
64     struct stat statBuffer;
65     size_t fileLen = 0;
66     size_t numSend = 0;
67     size_t divFlag = 0;
68     // Read the file length
69     if (stat(fileToSend, &statBuffer) < 0) {
70         // Handling reding file error
71         perror("Error: Can not read file\n");
72         exit(-1);
73     }
74     fileLen = statBuffer.st_size;
75     // 1.2 Calculate num of packets to be sent
76     divFlag = fileLen % MAX_PAYLOAD_SIZE;
77     numSend = divFlag > 0 ? (1 + (fileLen / MAX_PAYLOAD_SIZE))
78                  : fileLen / MAX_PAYLOAD_SIZE;
79     printf("File size is %lu, need %lu packets to send\n", fileLen, numSend);
80
81     // 2. send
82     unsigned int randSeq = 0;

```

```

83     char packet[MAX_MESSAGE_SIZE];
84     char msg[MAX_MESSAGE_SIZE];
85     struct SRHeader header;
86     memset(packet, 0, sizeof(packet));
87     // 2.1 send SYN, has random seq number
88     srand((unsigned)time(NULL));
89     randSeq = rand();
90     header.flag = SYN;
91     header.seq = randSeq;
92     header.len = 0;
93     header.crc = crc32(NULL, 0); // without payload;
94     memcpy(packet, (char *)&header, sizeof(header));
95     sendto(sockfd, (void *)packet, sizeof(header), 0,
96            (const struct sockaddr *)&recvAddr, sizeof(struct sockaddr_in));
97     logInfo(header, logfp);
98     printf("Send SYN to set up the connection\n");
99     // 2.2 wait for ACK and SYN packet
100    socklen_t sLen = sizeof(struct sockaddr_in);
101    struct sockaddr_in siOther;
102    struct SRHeader recvHeader;
103    struct timeval start, stop;
104    gettimeofday(&start, NULL);
105    gettimeofday(&stop, NULL);
106
107    // receive ack
108    while (1) {
109        if (recvfrom(sockfd, msg, MAX_MESSAGE_SIZE, MSG_DONTWAIT,
110                    (struct sockaddr *)&siOther, &sLen) > 0) {
111            memcpy(&recvHeader, msg, sizeof(recvHeader));
112            logInfo(recvHeader, logfp);
113            if (recvHeader.flag == ACK) {
114                printf("Connection set up \n");
115                break;
116            }
117        } else {
118            // resend SYN
119            sendto(sockfd, (void *)packet, sizeof(header), 0,
120                  (const struct sockaddr *)&recvAddr, sizeof(struct sockaddr_in));
121            logInfo(header, logfp);
122        }
123        // time out for ACK
124        gettimeofday(&stop, NULL);
125        if ((double)(stop.tv_sec - start.tv_sec) * 1000 +
126            (double)(stop.tv_usec - start.tv_usec) / 1000 >
127            5000) {
128            perror("Error: do not receive ACK for SYN within 5 s. Give up\n");
129            fclose(logfp);
130            exit(4);
131        }
132    }
133    // 2.3 Send data
134    char *buf = malloc(windowSize * MAX_MESSAGE_SIZE);
135    memset(buf, 0, windowSize * MAX_MESSAGE_SIZE);
136    FILE *fp = fopen(fileToSend, "r");
137    // open file
138    if (!fp) {
139        // Handling file open error
140        perror("Error: Can not read file\n");
141        free(buf);
142        fclose(logfp);
143        fclose(fp);
144        exit(-1);
145    }

```

```

146
147 unsigned int currentSeq = 0; // send 0 for initial
148 unsigned int bufNum = 0;
149 int endFlag = 1;
150 unsigned dupCnt = 0;
151 unsigned dupSeq = 0;
152 size_t payloadLen = MAX_PAYLOAD_SIZE;
153 gettimeofday(&start, NULL);
154 while (endFlag) {
155     // check if timeout
156     gettimeofday(&stop, NULL);
157     // if timeout, resend all buffer
158     if ((double)(stop.tv_sec - start.tv_sec) * 1000 +
159         (double)(stop.tv_usec - start.tv_usec) / 1000 >
160         400) {
161         printf(
162             "No ack within 400ms, resend all packets in buffer. Start from %u\n",
163             currentSeq);
164         // resend
165         for (unsigned int ii = 0; ii < bufNum; ii++) {
166             payloadLen = ((currentSeq + ii == numSend - 1) && (divFlag > 0))
167                 ? divFlag
168                 : MAX_PAYLOAD_SIZE;
169             sendto(sockfd, (void *) (buf + ii * sizeof(packet)),
170                 payloadLen + sizeof(header), 0,
171                 (const struct sockaddr *)&recvAddr, sizeof(struct sockaddr_in));
172             struct SRHeader tempHeader;
173             memcpy(&tempHeader, buf + ii * sizeof(packet), sizeof(header));
174             logInfo(tempHeader, logfp);
175         }
176         // reset time
177         gettimeofday(&start, NULL);
178     }
179     // send and buffer
180
181     if ((bufNum < windowSize) && ((bufNum + currentSeq) < numSend)) {
182         payloadLen = ((currentSeq + bufNum == numSend - 1) && (divFlag > 0))
183             ? divFlag
184             : MAX_PAYLOAD_SIZE;
185         // form a packet
186         fread(packet + sizeof(header), payloadLen, 1, fp);
187         header.flag = DATA;
188         header.seq = currentSeq + bufNum;
189         header.len = payloadLen;
190         header.crc = crc32(packet + sizeof(header), payloadLen);
191         memcpy(packet, (char *)(&header), sizeof(header));
192         // send
193         sendto(sockfd, (void *) packet, payloadLen + sizeof(header), 0,
194             (const struct sockaddr *)&recvAddr, sizeof(struct sockaddr_in));
195         logInfo(header, logfp);
196         // buffer the packet
197         memcpy(buf + bufNum * MAX_MESSAGE_SIZE, packet, sizeof(packet));
198         bufNum++;
199     }
200     // receive ACK
201     if (recvfrom(sockfd, msg, MAX_MESSAGE_SIZE, MSG_DONTWAIT,
202         (struct sockaddr *)&siOther, &sLen) != -1) {
203         memcpy(&recvHeader, msg, sizeof(recvHeader));
204         logInfo(recvHeader, logfp);
205         // duplicate ACK
206         if ((dupCnt == 0) || (recvHeader.seq != dupSeq)) {
207             dupSeq = recvHeader.seq;
208             dupCnt = 1;

```

```

209     } else if (recvHeader.seq == dupSeq) {
210         dupCnt++;
211         if (dupCnt >= 3) {
212             // resend it immediately
213             payloadLen = ((dupSeq == numSend - 1) && (divFlag > 0))
214                 ? divFlag
215                 : MAX_PAYLOAD_SIZE;
216             sendto(sockfd, (void *) (buf + (dupSeq - currentSeq) * sizeof(packet)),
217                 payloadLen + sizeof(header), 0,
218                 (const struct sockaddr *)&recvAddr,
219                 sizeof(struct sockaddr_in));
220             struct SRHeader tempHeader;
221             memcpy(&tempHeader, buf + (dupSeq - currentSeq) * sizeof(packet),
222                 sizeof(header));
223             logInfo(tempHeader, logfp);
224             // reset Cnt
225             dupCnt = 0;
226         }
227     }
228     // Advance window
229     if (recvHeader.flag == ACK) {
230         if (recvHeader.seq > currentSeq) {
231             // move buffer
232             if (recvHeader.seq > currentSeq + bufNum) {
233                 // clear buf
234                 memset(buf, 0, windowSize * MAX_MESSAGE_SIZE);
235                 bufNum = 0;
236             } else {
237                 char *temp = malloc(windowSize * MAX_PAYLOAD_SIZE);
238                 memcpy(temp, buf + (recvHeader.seq - currentSeq) * MAX_MESSAGE_SIZE,
239                     (bufNum + currentSeq - recvHeader.seq) * MAX_MESSAGE_SIZE);
240                 memset(buf, 0, windowSize * MAX_MESSAGE_SIZE);
241                 memcpy(buf, temp, windowSize * MAX_MESSAGE_SIZE);
242                 free(temp);
243                 bufNum = bufNum + currentSeq - recvHeader.seq;
244             }
245             // advance
246             currentSeq = recvHeader.seq;
247             if (currentSeq > (numSend - 1))
248                 endFlag = 0;
249         }
250     }
251     // reset time
252     gettimeofday(&start, NULL);
253 }
254 }
255 // Send FIN
256 header.flag = FIN;
257 header.seq = randSeq;
258 header.len = 0;
259 header.crc = crc32(NULL, 0); // without payload;
260 memcpy(packet, (char *)&header, sizeof(header));
261 sendto(sockfd, (void *) packet, sizeof(header), 0,
262     (const struct sockaddr *)&recvAddr, sizeof(struct sockaddr_in));
263 logInfo(header, logfp);
264 printf("Transmission finished: Send FIN");
265 // receive ack for FIN
266 while (1) {
267     if (recvfrom(sockfd, msg, MAX_MESSAGE_SIZE, MSG_DONTWAIT,
268         (struct sockaddr *)&siOther, &sLen) != -1) {
269         memcpy(&recvHeader, msg, sizeof(recvHeader));
270         logInfo(recvHeader, logfp);
271         if (recvHeader.flag == ACK) {

```

```

272     printf("Receive ACK for FIN, done\n");
273     break;
274 }
275 } else {
276     // resend FIN
277     sendto(sockfd, (void *)packet, sizeof(header), 0,
278             (const struct sockaddr *)&recvAddr, sizeof(struct sockaddr_in));
279     logInfo(header, logfp);
280 }
281 // time out for ACK
282 gettimeofday(&stop, NULL);
283 if ((double)(stop.tv_sec - start.tv_sec) * 1000 +
284     (double)(stop.tv_usec - start.tv_usec) / 1000 >
285     400) {
286     perror("Error: do not receive ACK for FIN within 400 ms. Give up\n");
287     fclose(fp);
288     fclose(logfp);
289     free(buf);
290     exit(4);
291 }
292 }
293 fclose(fp);
294 fclose(logfp);
295 free(buf);
296 return 0;
297 }
298
299 int runReceiver(int port, unsigned int windowSize, char *recvDir,
300               char *logFile) {
301     FILE *logfp = fopen(logFile, "w+");
302     // Set up connection
303     struct sockaddr_in siMe;
304     int sockfd = socket(AF_INET, SOCK_DGRAM, IPPROTO_UDP);
305     if (sockfd == -1) {
306         fputs("socket creation failed!", stderr);
307         exit(2); // error, exit with 2
308     }
309     memset((char *)&siMe, 0, sizeof(struct sockaddr_in));
310     siMe.sin_family = AF_INET;
311     siMe.sin_port = htons(port);
312     siMe.sin_addr.s_addr = htonl(INADDR_ANY);
313     bind(sockfd, (struct sockaddr *)&siMe, sizeof(struct sockaddr_in));
314
315     socklen_t sLen = sizeof(struct sockaddr_in);
316     struct sockaddr_in siOther;
317     struct timeval start, stop;
318     int connectCnt = 0;
319     while (1) {
320         // 1. Ready for receiving SYN
321         char msg[MAX_MESSAGE_SIZE];
322         memset(msg, 0, MAX_MESSAGE_SIZE);
323         char ACKMsg[MAX_MESSAGE_SIZE];
324         memset(ACKMsg, 0, sizeof(ACKMsg));
325         unsigned SYNFINSeq = 0;
326         struct SRHeader recvHeader;
327         struct SRHeader ACKHeader;
328
329         gettimeofday(&start, NULL);
330         gettimeofday(&stop, NULL);
331         while (1) {
332             if (recvfrom(sockfd, msg, MAX_MESSAGE_SIZE, MSG_DONTWAIT,
333                         (struct sockaddr *)&siOther, &sLen) > 0) {
334                 memcpy(&recvHeader, msg, sizeof(recvHeader));

```

```

335     logInfo(recvHeader, logfp);
336     if (recvHeader.flag == SYN) {
337         SYNFINSeq = recvHeader.seq;
338         // ACK for SYN
339         ACKHeader.flag = ACK;
340         ACKHeader.seq = SYNFINSeq;
341         ACKHeader.len = 0;
342         ACKHeader.crc = crc32(NULL, 0);
343         memcpy(ACKMsg, (char *)&ACKHeader, sizeof(ACKHeader));
344         sendto(sockfd, (void *)ACKMsg, sizeof(ACKHeader), 0,
345             (const struct sockaddr *)&siOther, sizeof(struct sockaddr_in));
346         logInfo(ACKHeader, logfp);
347         break;
348     }
349 }
350 gettimeofday(&stop, NULL);
351 if (((double)(stop.tv_sec - start.tv_sec) * 1000 +
352     (double)(stop.tv_usec - start.tv_usec) / 1000 >
353     5000) {
354     printf("I do not receive SYN within 5 s. Give up\n");
355     return 0;
356 }
357 }
358 // 2. Recive data
359
360 // 2.1 open file
361 // create file name
362 char fileName[30];
363 char name[13];
364 sprintf(name, "file_%d.txt", connectCnt);
365 strcpy(fileName, recvDir);
366 strcat(fileName, name);
367 FILE *fp = fopen(fileName, "w");
368 unsigned int currentACKSeq = 0;
369 char *buf = malloc(windowSize * (MAX_PAYLOAD_SIZE + 8)); // valid. len
370 memset(buf, 0, windowSize * (MAX_PAYLOAD_SIZE + 8));
371 int FINflag = 1; // 0 for receive FIN flag
372
373 // 2.2 receive data
374 printf("Start to receive file\n");
375 while (FINflag) {
376     // receive data
377     if (recvfrom(sockfd, msg, MAX_MESSAGE_SIZE, MSG_DONTWAIT,
378         (struct sockaddr *)&siOther, &sLen) != -1) {
379         memcpy(&recvHeader, msg, sizeof(recvHeader));
380         logInfo(recvHeader, logfp);
381         // If it's data
382         if (recvHeader.flag == DATA) {
383             if (recvHeader.crc ==
384                 (crc32(msg + sizeof(recvHeader), recvHeader.len))) {
385                 // check the condition for window operation
386                 if (recvHeader.seq >= currentACKSeq + windowSize) {
387                     continue;
388                 } else if (recvHeader.seq < currentACKSeq) {
389                     // send ACK currentACKSeq
390                     ACKHeader.flag = ACK;
391                     ACKHeader.seq = currentACKSeq;
392                     ACKHeader.len = 0;
393                     ACKHeader.crc = crc32(NULL, 0);
394                     memcpy(ACKMsg, (char *)&ACKHeader, sizeof(ACKHeader));
395                     sendto(sockfd, (void *)ACKMsg, sizeof(ACKHeader), 0,
396                         (const struct sockaddr *)&siOther,
397                         sizeof(struct sockaddr_in));

```

```

398     logInfo(ACKHeader, logfp);
399 } else if ((recvHeader.seq > currentACKSeq) &&
400           (recvHeader.seq < currentACKSeq + windowSize)) {
401     // buffer it
402     unsigned int temp = 1;
403     memcpy(buf + (recvHeader.seq - currentACKSeq) *
404            (MAX_PAYLOAD_SIZE + 8),
405            (char *)&temp, sizeof(temp));
406     memcpy(buf +
407            (recvHeader.seq - currentACKSeq) *
408            (MAX_PAYLOAD_SIZE + 8) +
409            4,
410            (char *)&recvHeader.len, sizeof(temp));
411     memcpy(buf +
412            + (recvHeader.seq - currentACKSeq) *
413            (MAX_PAYLOAD_SIZE + 8) +
414            8,
415            msg + sizeof(recvHeader), recvHeader.len);
416     // send current ACK
417     ACKHeader.flag = ACK;
418     ACKHeader.seq = currentACKSeq;
419     ACKHeader.len = 0;
420     ACKHeader.crc = crc32(NULL, 0);
421     memcpy(ACKMsg, (char *)&ACKHeader, sizeof(ACKHeader));
422     sendto(sockfd, (void *)ACKMsg, sizeof(ACKHeader), 0,
423            (const struct sockaddr *)&siOther,
424            sizeof(struct sockaddr_in));
425     logInfo(ACKHeader, logfp);
426 } else if (recvHeader.seq == currentACKSeq) {
427     // calculate the highest available packet
428     unsigned int numWrite = 1;
429     for (numWrite = 1; numWrite < windowSize; numWrite++) {
430         unsigned int temp;
431         memcpy(&temp, buf + numWrite * (8 + MAX_PAYLOAD_SIZE),
432                sizeof(temp));
433         if (temp == 0)
434             break;
435     }
436     fwrite(msg + sizeof(recvHeader), 1, recvHeader.len, fp);
437     for (unsigned int ii = 1; ii < numWrite; ii++) {
438         unsigned int tempLen;
439         memcpy(&tempLen, buf + ii * (8 + MAX_PAYLOAD_SIZE) + 4,
440                sizeof(tempLen));
441         fwrite(buf + ii * (8 + MAX_PAYLOAD_SIZE) + 8, 1, tempLen, fp);
442     }
443     // clear buff
444     memset(buf, 0, windowSize * (MAX_PAYLOAD_SIZE + 4));
445     // ACK
446     currentACKSeq = currentACKSeq + numWrite;
447     // send current ACK
448     ACKHeader.flag = ACK;
449     ACKHeader.seq = currentACKSeq;
450     ACKHeader.len = 0;
451     ACKHeader.crc = crc32(NULL, 0);
452     memcpy(ACKMsg, (char *)&ACKHeader, sizeof(ACKHeader));
453     sendto(sockfd, (void *)ACKMsg, sizeof(ACKHeader), 0,
454            (const struct sockaddr *)&siOther,
455            sizeof(struct sockaddr_in));
456     logInfo(ACKHeader, logfp);
457 }
458 }
459 } else if (recvHeader.flag == FIN) {
460     FINflag = 0;

```



```

461         // send ACK
462         ACKHeader.flag = ACK;
463         ACKHeader.seq = SYNFINSeq;
464         ACKHeader.len = 0;
465         ACKHeader.crc = crc32(NULL, 0);
466         memcpy(ACKMsg, (char *)&ACKHeader, sizeof(ACKHeader));
467         sendto(sockfd, (void *)ACKMsg, sizeof(ACKHeader), 0,
468                 (const struct sockaddr *)&siOther, sizeof(struct sockaddr_in));
469         logInfo(ACKHeader, logfp);
470         // end
471         fclose(fp);
472         free(buf);
473         connectCnt++;
474         fclose(logfp);
475     }
476 }
477 }
478 }
479 fclose(logfp);
480 return 0;
481 }
482
483 int main(int argc, char **argv) {
484     // Sender mode
485     if (strcmp(argv[1], "-s") == 0) {
486         // Error for number of arguments
487         if (argc != 8) {
488             perror(
489                 "./sr -s <receiver's IP> <receiver's port> <sender's port> <window "
490                 "size> <file to send> <log file>");
491             return 1; // error: exit with 1
492         }
493         char *recvIP = argv[2];
494         int recvPort = atoi(argv[3]);
495         int senderPort = atoi(argv[4]);
496         unsigned int windowSize = atoi(argv[5]);
497         char *fileToSend = argv[6];
498         char *logFile = argv[7];
499         runSender(recvIP, recvPort, senderPort, windowSize, fileToSend, logFile);
500         return 0;
501     }
502
503     // Receiver mode
504     else if (strcmp(argv[1], "-r") == 0) {
505         // Error for number of arguments
506         if (argc != 6) {
507             perror("./sr -r <port> <window size> <recv dir> <log file>");
508             return 1; // error: exit with 1
509         }
510         // int runReceiver(int port, unsigned int windowSize, char *recvDir,
511         //                 char *logFile)
512         int port = atoi(argv[2]);
513         unsigned int windowSize = atoi(argv[3]);
514         char *recvDir = argv[4];
515         char *logFile = argv[5];
516         runReceiver(port, windowSize, recvDir, logFile);
517         return 0;
518     } else {
519         perror("./sr -s <receiver's IP> <receiver's port> <sender's port> <window "
520             "size> <file to send> <log file>\n ./sr -r <port> <window size> "
521             "<recv dir> <log file>\n");
522         return 1;
523     }

```

```

524     return 0;
525 }

```

4.2.4 sr.c in 3.3

```

1  #include "../SRHeader.h"
2  #include "../crc32.h"
3  #include <arpa/inet.h>
4  #include <netdb.h>
5  #include <stdint.h>
6  #include <stdio.h>
7  #include <stdlib.h>
8  #include <string.h>
9  #include <sys/socket.h>
10 #include <sys/stat.h> // struct stat
11 #include <sys/time.h>
12 #include <time.h>
13 #include <unistd.h>
14 // The size of packet would is len(header) + len(payload) = 4*4 + 1456 = 1472
15 #define MAX_MESSAGE_SIZE 1472
16
17 int logInfo(struct SRHeader header, FILE *fp) {
18     if (header.flag == SYN) {
19         fprintf(fp, "%s %u %u %u\n", "SYN", header.seq, header.len, header.crc);
20         printf("%s %u %u %u\n", "SYN", header.seq, header.len, header.crc);
21     } else if (header.flag == FIN) {
22         fprintf(fp, "%s %u %u %u\n", "FIN", header.seq, header.len, header.crc);
23         printf("%s %u %u %u\n", "FIN", header.seq, header.len, header.crc);
24     } else if (header.flag == DATA) {
25         fprintf(fp, "%s %u %u %u\n", "DATA", header.seq, header.len, header.crc);
26         printf("%s %u %u %u\n", "DATA", header.seq, header.len, header.crc);
27     } else if (header.flag == ACK) {
28         fprintf(fp, "%s %u %u %u\n", "ACK", header.seq, header.len, header.crc);
29         printf("%s %u %u %u\n", "ACK", header.seq, header.len, header.crc);
30     } else if (header.flag == NACK) {
31         fprintf(fp, "%s %u %u %u\n", "NACK", header.seq, header.len, header.crc);
32         printf("%s %u %u %u\n", "NACK", header.seq, header.len, header.crc);
33     }
34     return 0;
35 }
36
37 int runSender(char *recvIP, int recvPort, int senderPort,
38               unsigned int windowSize, char *fileToSend, char *logFile) {
39     FILE *logfp = fopen(logFile, "w+");
40     // create UDP socket.
41     int sockfd = socket(AF_INET, SOCK_DGRAM, IPPROTO_UDP); // use IPPROTO_UDP
42     if (sockfd < 0) {
43         perror("socket creation failed");
44         exit(1);
45     }
46
47     // below is the same as TCP socket.
48     struct sockaddr_in myAddr;
49     memset(&myAddr, 0, sizeof(myAddr));
50     myAddr.sin_family = AF_INET;
51     myAddr.sin_addr.s_addr = INADDR_ANY;
52     // bind to a specific port
53     myAddr.sin_port = htons(senderPort);
54     bind(sockfd, (struct sockaddr *)&myAddr, sizeof(myAddr));
55
56     struct sockaddr_in recvAddr;
57     struct hostent *host = gethostbyname(recvIP);
58     memcpy(&(recvAddr.sin_addr), host->h_addr, host->h_length);

```

```

59     recvAddr.sin_family = AF_INET;
60     recvAddr.sin_port = htons(recvPort);
61
62     // 1. The num of packets to be sent
63     // 1.1 Read file length
64     struct stat statBuffer;
65     size_t fileLen = 0;
66     size_t numSend = 0;
67     size_t divFlag = 0;
68     // Read the file length
69     if (stat(fileToSend, &statBuffer) < 0) {
70         // Handling reading file error
71         perror("Error: Can not read file\n");
72         exit(-1);
73     }
74     fileLen = statBuffer.st_size;
75     // 1.2 Calculate num of packets to be sent
76     divFlag = fileLen % MAX_PAYLOAD_SIZE;
77     numSend = divFlag > 0 ? (1 + (fileLen / MAX_PAYLOAD_SIZE))
78                 : fileLen / MAX_PAYLOAD_SIZE;
79     printf("File size is %lu, need %lu packets to send\n", fileLen, numSend);
80
81     // 2. send
82     unsigned int randSeq = 0;
83     char packet[MAX_MESSAGE_SIZE];
84     char msg[MAX_MESSAGE_SIZE];
85     struct SRHeader header;
86     memset(packet, 0, sizeof(packet));
87     // 2.1 send SYN, has random seq number
88     srand((unsigned)time(NULL));
89     randSeq = rand();
90     header.flag = SYN;
91     header.seq = randSeq;
92     header.len = 0;
93     header.crc = crc32(NULL, 0); // without payload;
94     memcpy(packet, (char *)&header, sizeof(header));
95     sendto(sockfd, (void *)packet, sizeof(header), 0,
96            (const struct sockaddr *)&recvAddr, sizeof(struct sockaddr_in));
97     logInfo(header, logfp);
98     printf("Send SYN to set up the connection\n");
99     // 2.2 wait for ACK and SYN packet
100    socklen_t sLen = sizeof(struct sockaddr_in);
101    struct sockaddr_in siOther;
102    struct SRHeader recvHeader;
103    struct timeval start, stop;
104    gettimeofday(&start, NULL);
105    gettimeofday(&stop, NULL);
106
107    // receive ack
108    while (1) {
109        if (recvfrom(sockfd, msg, MAX_MESSAGE_SIZE, MSG_DONTWAIT,
110                    (struct sockaddr *)&siOther, &sLen) > 0) {
111            memcpy(&recvHeader, msg, sizeof(recvHeader));
112            logInfo(recvHeader, logfp);
113            if (recvHeader.flag == ACK) {
114                printf("Connection set up \n");
115                break;
116            }
117        } else {
118            // resend SYN
119            sendto(sockfd, (void *)packet, sizeof(header), 0,
120                  (const struct sockaddr *)&recvAddr, sizeof(struct sockaddr_in));
121            logInfo(header, logfp);

```

```

122     }
123     // time out for ACK
124     gettimeofday(&stop, NULL);
125     if ((double)(stop.tv_sec - start.tv_sec) * 1000 +
126         (double)(stop.tv_usec - start.tv_usec) / 1000 >
127         5000) {
128         perror("Error: do not receive ACK for SYN within 5 s. Give up\n");
129         fclose(logfp);
130         exit(4);
131     }
132 }
133 // 2.3 Send data
134 char *buf = malloc(windowSize * MAX_MESSAGE_SIZE);
135 memset(buf, 0, windowSize * MAX_MESSAGE_SIZE);
136 FILE *fp = fopen(fileToSend, "r");
137 // open file
138 if (!fp) {
139     // Handling file open error
140     perror("Error: Can not read file\n");
141     free(buf);
142     fclose(logfp);
143     fclose(fp);
144     exit(-1);
145 }
146
147 unsigned int currentSeq = 0; // send 0 for initial
148 unsigned int bufNum = 0;
149 int endFlag = 1;
150 size_t payloadLen = MAX_PAYLOAD_SIZE;
151 gettimeofday(&start, NULL);
152 while (endFlag) {
153     // check if timeout
154     gettimeofday(&stop, NULL);
155     // if timeout, resend all buffer
156     if ((double)(stop.tv_sec - start.tv_sec) * 1000 +
157         (double)(stop.tv_usec - start.tv_usec) / 1000 >
158         400) {
159         printf(
160             "No ack within 400ms, resend all packets in buffer. Start from %u\n",
161             currentSeq);
162         // resend
163         for (unsigned int ii = 0; ii < bufNum; ii++) {
164             payloadLen = ((currentSeq + ii == numSend - 1) && (divFlag > 0))
165                 ? divFlag
166                 : MAX_PAYLOAD_SIZE;
167             sendto(sockfd, (void *) (buf + ii * sizeof(packet)),
168                 payloadLen + sizeof(header), 0,
169                 (const struct sockaddr *)&recvAddr, sizeof(struct sockaddr_in));
170             struct SRHeader tempHeader;
171             memcpy(&tempHeader, buf + ii * sizeof(packet), sizeof(header));
172             logInfo(tempHeader, logfp);
173         }
174         // reset time
175         gettimeofday(&start, NULL);
176     }
177     // send and buffer
178     if ((bufNum < windowSize) && ((bufNum + currentSeq) < numSend)) {
179         payloadLen = ((currentSeq + bufNum == numSend - 1) && (divFlag > 0))
180             ? divFlag
181             : MAX_PAYLOAD_SIZE;
182         // form a packet
183         fread(packet + sizeof(header), payloadLen, 1, fp);
184         header.flag = DATA;

```

```

185     header.seq = currentSeq + bufNum;
186     header.len = payloadLen;
187     header.crc = crc32(packet + sizeof(header), payloadLen);
188     memcpy(packet, (char *)&header, sizeof(header));
189     // send
190     sendto(sockfd, (void *)packet, payloadLen + sizeof(header), 0,
191             (const struct sockaddr *)&recvAddr, sizeof(struct sockaddr_in));
192     logInfo(header, logfp);
193     // buffer the packet
194     memcpy(buf + bufNum * MAX_MESSAGE_SIZE, packet, sizeof(packet));
195     bufNum++;
196 }
197 // receive ACK
198 if (recvfrom(sockfd, msg, MAX_MESSAGE_SIZE, MSG_DONTWAIT,
199             (struct sockaddr *)&siOther, &sLen) != -1) {
200     memcpy(&recvHeader, msg, sizeof(recvHeader));
201     logInfo(recvHeader, logfp);
202     // Advance window
203     if (recvHeader.flag == ACK) {
204         if (recvHeader.seq > currentSeq) {
205             // move buffer
206             if (recvHeader.seq > currentSeq + bufNum) {
207                 // clear buf
208                 memset(buf, 0, windowSize * MAX_MESSAGE_SIZE);
209                 bufNum = 0;
210             } else {
211                 char *temp = malloc(windowSize * MAX_PAYLOAD_SIZE);
212                 memcpy(temp, buf + (recvHeader.seq - currentSeq) * MAX_MESSAGE_SIZE,
213                        (bufNum + currentSeq - recvHeader.seq) * MAX_MESSAGE_SIZE);
214                 memset(buf, 0, windowSize * MAX_MESSAGE_SIZE);
215                 memcpy(buf, temp, windowSize * MAX_MESSAGE_SIZE);
216                 free(temp);
217                 bufNum = bufNum + currentSeq - recvHeader.seq;
218             }
219             // advance
220             currentSeq = recvHeader.seq;
221             if (currentSeq > (numSend - 1))
222                 endFlag = 0;
223         }
224         // resend if NACK
225     } else if (recvHeader.flag == NACK) {
226         payloadLen = ((recvHeader.seq == numSend - 1) && (divFlag > 0))
227                     ? divFlag
228                     : MAX_PAYLOAD_SIZE;
229         sendto(sockfd,
230                (void *) (buf + (recvHeader.seq - currentSeq) * sizeof(packet)),
231                payloadLen + sizeof(header), 0,
232                (const struct sockaddr *)&recvAddr, sizeof(struct sockaddr_in));
233         struct SRHeader tempHeader;
234         memcpy(&tempHeader,
235                buf + (recvHeader.seq - currentSeq) * sizeof(packet),
236                sizeof(header));
237         logInfo(tempHeader, logfp);
238     }
239     // reset time
240     gettimeofday(&start, NULL);
241 }
242 }
243 // Send FIN
244 header.flag = FIN;
245 header.seq = randSeq;
246 header.len = 0;
247 header.crc = crc32(NULL, 0); // without payload;

```

```

248 memcpy(packet, (char *)&header, sizeof(header));
249 sendto(sockfd, (void *)packet, sizeof(header), 0,
250         (const struct sockaddr *)&recvAddr, sizeof(struct sockaddr_in));
251 logInfo(header, logfp);
252 printf("Transmission finished: Send FIN");
253 // receive ack for FIN
254 while (1) {
255     if (recvfrom(sockfd, msg, MAX_MESSAGE_SIZE, MSG_DONTWAIT,
256                 (struct sockaddr *)&siOther, &sLen) != -1) {
257         memcpy(&recvHeader, msg, sizeof(recvHeader));
258         logInfo(recvHeader, logfp);
259         if (recvHeader.flag == ACK) {
260             printf("Receive ACK for FIN, done\n");
261             break;
262         }
263     } else {
264         // resend FIN
265         sendto(sockfd, (void *)packet, sizeof(header), 0,
266               (const struct sockaddr *)&recvAddr, sizeof(struct sockaddr_in));
267         logInfo(header, logfp);
268     }
269     // time out for ACK
270     gettimeofday(&stop, NULL);
271     if ((double)(stop.tv_sec - start.tv_sec) * 1000 +
272         (double)(stop.tv_usec - start.tv_usec) / 1000 >
273         400) {
274         perror("Error: do not receive ACK for FIN within 400 ms. Give up\n");
275         fclose(fp);
276         fclose(logfp);
277         free(buf);
278         exit(4);
279     }
280 }
281 fclose(fp);
282 fclose(logfp);
283 free(buf);
284 return 0;
285 }
286
287 int runReceiver(int port, unsigned int windowSize, char *recvDir,
288                char *logFile) {
289     FILE *logfp = fopen(logFile, "w+");
290     // Set up connection
291     struct sockaddr_in siMe;
292     int sockfd = socket(AF_INET, SOCK_DGRAM, IPPROTO_UDP);
293     if (sockfd == -1) {
294         fputs("socket creation failed!", stderr);
295         exit(2); // error, exit with 2
296     }
297     memset((char *)&siMe, 0, sizeof(struct sockaddr_in));
298     siMe.sin_family = AF_INET;
299     siMe.sin_port = htons(port);
300     siMe.sin_addr.s_addr = htonl(INADDR_ANY);
301     bind(sockfd, (struct sockaddr *)&siMe, sizeof(struct sockaddr_in));
302
303     socklen_t sLen = sizeof(struct sockaddr_in);
304     struct sockaddr_in siOther;
305     struct timeval start, stop;
306     int connectCnt = 0;
307     while (1) {
308         // 1. Ready for receiving SYN
309         char msg[MAX_MESSAGE_SIZE];
310         memset(msg, 0, MAX_MESSAGE_SIZE);

```

```

311     char ACKMsg[MAX_MESSAGE_SIZE];
312     memset(ACKMsg, 0, sizeof(ACKMsg));
313     unsigned SYNFINSeq = 0;
314     struct SRHeader recvHeader;
315     struct SRHeader ACKHeader;
316
317     gettimeofday(&start, NULL);
318     gettimeofday(&stop, NULL);
319     while (1) {
320         if (recvfrom(sockfd, msg, MAX_MESSAGE_SIZE, MSG_DONTWAIT,
321             (struct sockaddr *)&siOther, &sLen) > 0) {
322             memcpy(&recvHeader, msg, sizeof(recvHeader));
323             logInfo(recvHeader, logfp);
324             if (recvHeader.flag == SYN) {
325                 SYNFINSeq = recvHeader.seq;
326                 // ACK for SYN
327                 ACKHeader.flag = ACK;
328                 ACKHeader.seq = SYNFINSeq;
329                 ACKHeader.len = 0;
330                 ACKHeader.crc = crc32(NULL, 0);
331                 memcpy(ACKMsg, (char *)&ACKHeader, sizeof(ACKHeader));
332                 sendto(sockfd, (void *)ACKMsg, sizeof(ACKHeader), 0,
333                     (const struct sockaddr *)&siOther, sizeof(struct sockaddr_in));
334                 logInfo(ACKHeader, logfp);
335                 break;
336             }
337         }
338         gettimeofday(&stop, NULL);
339         if ((double)(stop.tv_sec - start.tv_sec) * 1000 +
340             (double)(stop.tv_usec - start.tv_usec) / 1000 >
341             5000) {
342             printf("I do not receive SYN within 5 s. Give up\n");
343             return 0;
344         }
345     }
346     // 2. Recive data
347
348     // 2.1 open file
349     // create file name
350     char fileName[30];
351     char name[13];
352     sprintf(name, "file_%d.txt", connectCnt);
353     strcpy(fileName, recvDir);
354     strcat(fileName, name);
355     FILE *fp = fopen(fileName, "w");
356     unsigned int currentACKSeq = 0;
357     char *buf = malloc(windowSize * (MAX_PAYLOAD_SIZE + 8)); // valid. len
358     memset(buf, 0, windowSize * (MAX_PAYLOAD_SIZE + 8));
359     int FINflag = 1; // 0 for receive FIN flag
360     int advanceFlag = 0;
361     // 2.2 receive data
362     printf("Start to receive file\n");
363     while (FINflag) {
364         advanceFlag = 0;
365         // receive data
366         if (recvfrom(sockfd, msg, MAX_MESSAGE_SIZE, MSG_DONTWAIT,
367             (struct sockaddr *)&siOther, &sLen) != -1) {
368             memcpy(&recvHeader, msg, sizeof(recvHeader));
369             logInfo(recvHeader, logfp);
370             // If it's data
371             if (recvHeader.flag == DATA) {
372                 if (recvHeader.crc ==
373                     (crc32(msg + sizeof(recvHeader), recvHeader.len))) {

```

```

374 // check the condition for window operation
375 if (recvHeader.seq >= currentACKSeq + windowSize) {
376     continue;
377 } else if (recvHeader.seq < currentACKSeq) {
378     // send ACK currentACKSeq
379     ACKHeader.flag = ACK;
380     ACKHeader.seq = currentACKSeq;
381     ACKHeader.len = 0;
382     ACKHeader.crc = crc32(NULL, 0);
383     memcpy(ACKMsg, (char *)&ACKHeader, sizeof(ACKHeader));
384     sendto(sockfd, (void *)ACKMsg, sizeof(ACKHeader), 0,
385             (const struct sockaddr *)&siOther,
386             sizeof(struct sockaddr_in));
387     logInfo(ACKHeader, logfp);
388 } else if ((recvHeader.seq > currentACKSeq) &&
389            (recvHeader.seq < currentACKSeq + windowSize)) {
390     // buffer it
391     unsigned int temp = 1;
392     memcpy(buf + (recvHeader.seq - currentACKSeq) *
393            (MAX_PAYLOAD_SIZE + 8),
394            (char *)&temp, sizeof(temp));
395     memcpy(buf +
396            (recvHeader.seq - currentACKSeq) *
397            (MAX_PAYLOAD_SIZE + 8) +
398            4,
399            (char *)&recvHeader.len, sizeof(temp));
400     memcpy(buf +
401            + (recvHeader.seq - currentACKSeq) *
402            (MAX_PAYLOAD_SIZE + 8) +
403            8,
404            msg + sizeof(recvHeader), recvHeader.len);
405     // send current ACK
406     ACKHeader.flag = ACK;
407     ACKHeader.seq = currentACKSeq;
408     ACKHeader.len = 0;
409     ACKHeader.crc = crc32(NULL, 0);
410     memcpy(ACKMsg, (char *)&ACKHeader, sizeof(ACKHeader));
411     sendto(sockfd, (void *)ACKMsg, sizeof(ACKHeader), 0,
412            (const struct sockaddr *)&siOther,
413            sizeof(struct sockaddr_in));
414     logInfo(ACKHeader, logfp);
415 } else if (recvHeader.seq == currentACKSeq) {
416     advanceFlag = 1;
417     // calculate the highest available packet
418     unsigned int numWrite = 1;
419     for (numWrite = 1; numWrite < windowSize; numWrite++) {
420         unsigned int temp;
421         memcpy(&temp, buf + numWrite * (8 + MAX_PAYLOAD_SIZE),
422                sizeof(temp));
423         if (temp == 0)
424             break;
425     }
426     fwrite(msg + sizeof(recvHeader), 1, recvHeader.len, fp);
427     for (unsigned int ii = 1; ii < numWrite; ii++) {
428         unsigned templen;
429         memcpy(&templen, buf + ii * (8 + MAX_PAYLOAD_SIZE) + 4,
430                sizeof(templen));
431         fwrite(buf + ii * (8 + MAX_PAYLOAD_SIZE) + 8, 1, templen, fp);
432     }
433     // clear buff
434     memset(buf, 0, windowSize * (MAX_PAYLOAD_SIZE + 4));
435     // ACK
436     currentACKSeq = currentACKSeq + numWrite;

```



```

437         // send current ACK
438         ACKHeader.flag = ACK;
439         ACKHeader.seq = currentACKSeq;
440         ACKHeader.len = 0;
441         ACKHeader.crc = crc32(NULL, 0);
442         memcpy(ACKMsg, (char *)&ACKHeader, sizeof(ACKHeader));
443         sendto(sockfd, (void *)ACKMsg, sizeof(ACKHeader), 0,
444             (const struct sockaddr *)&siOther,
445             sizeof(struct sockaddr_in));
446         logInfo(ACKHeader, logfp);
447     }
448 }
449 // check for gap and send NACK
450 unsigned int continueCnt = 1;
451 unsigned int NACKFlag = 0;
452 for (continueCnt = 0; continueCnt < windowSize; continueCnt++) {
453     unsigned int temp;
454     memcpy(&temp, buf + continueCnt * (8 + MAX_PAYLOAD_SIZE),
455         sizeof(temp));
456     if (temp == 0)
457         break;
458 }
459 if (continueCnt < windowSize) {
460     if (continueCnt == 0) {
461         if (!advanceFlag)
462             NACKFlag = 1;
463     } else
464         NACKFlag = 1;
465     if (NACKFlag) {
466         ACKHeader.flag = NACK;
467         ACKHeader.seq = currentACKSeq + continueCnt;
468         ACKHeader.len = 0;
469         ACKHeader.crc = crc32(NULL, 0);
470         memcpy(ACKMsg, (char *)&ACKHeader, sizeof(ACKHeader));
471         sendto(sockfd, (void *)ACKMsg, sizeof(ACKHeader), 0,
472             (const struct sockaddr *)&siOther,
473             sizeof(struct sockaddr_in));
474         logInfo(ACKHeader, logfp);
475     }
476 }
477 } else if (recvHeader.flag == FIN) {
478     FINflag = 0;
479     // send ACK
480     ACKHeader.flag = ACK;
481     ACKHeader.seq = SYNFINSeq;
482     ACKHeader.len = 0;
483     ACKHeader.crc = crc32(NULL, 0);
484     memcpy(ACKMsg, (char *)&ACKHeader, sizeof(ACKHeader));
485     sendto(sockfd, (void *)ACKMsg, sizeof(ACKHeader), 0,
486         (const struct sockaddr *)&siOther, sizeof(struct sockaddr_in));
487     logInfo(ACKHeader, logfp);
488     // end
489     fclose(fp);
490     free(buf);
491     connectCnt++;
492     fclose(logfp);
493 }
494 }
495 }
496 }
497 fclose(logfp);
498 return 0;
499 }

```

```

500
501 int main(int argc, char **argv) {
502     // Sender mode
503     if (strcmp(argv[1], "-s") == 0) {
504         // Error for number of arguments
505         if (argc != 8) {
506             perror(
507                 "./sr -s <receiver's IP> <receiver's port> <sender's port> <window "
508                 "size> <file to send> <log file>");
509             return 1; // error: exit with 1
510         }
511         char *recvIP = argv[2];
512         int recvPort = atoi(argv[3]);
513         int senderPort = atoi(argv[4]);
514         unsigned int windowSize = atoi(argv[5]);
515         char *fileToSend = argv[6];
516         char *logFile = argv[7];
517         runSender(recvIP, recvPort, senderPort, windowSize, fileToSend, logFile);
518         return 0;
519     }
520
521     // Receiver mode
522     else if (strcmp(argv[1], "-r") == 0) {
523         // Error for number of arguments
524         if (argc != 6) {
525             perror("./sr -r <port> <window size> <recv dir> <log file>");
526             return 1; // error: exit with 1
527         }
528         // int runReceiver(int port, unsigned int windowSize, char *recvDir,
529         //                 char *logFile)
530         int port = atoi(argv[2]);
531         unsigned int windowSize = atoi(argv[3]);
532         char *recvDir = argv[4];
533         char *logFile = argv[5];
534         runReceiver(port, windowSize, recvDir, logFile);
535         return 0;
536     } else {
537         perror("./sr -s <receiver's IP> <receiver's port> <sender's port> <window "
538             "size> <file to send> <log file>\n ./sr -r <port> <window size> "
539             "<recv dir> <log file>\n");
540         return 1;
541     }
542     return 0;
543 }

```